

Lightweight Application-Level Crash Consistency on Transactional Flash Storage

Changwoo Min, Woon-Hak Kang[†], Taesoo Kim, Sang-Won Lee[†], Young Ik Eom[†]




Georgia Institute of Technology
[†]Sungkyunkwan University




Application's data is not consistent after random failures

Mobile Phone



**Hang
while booting**

Bank Account



Financial loss

Application's data is not consistent after random failures

Mobile Phone



Bank Account



Power Outage
Hardware Errors
Software Panics (OS, Device Driver)

**Hang
while booting**

Financial loss

Example: inserting records in two databases

Application

```
write(/db1, "new");  
write(/db2, "new");
```

Example: inserting records in two databases

Application


```
write(/db1, "new");  
write(/db2, "new");
```

File System

For each database

- Allocates new data block (bitmap)
- Fills data block with user data (data)
- Sets location of data block (inode)

Storage Device



Example: inserting records in two databases

Application


```
write(/db1, "new");  
write(/db2, "new");
```

File System

For each database

- Allocates new data block (bitmap)
- Fills data block with user data (data)
- Sets location of data block (inode)

Storage Device



Example: inserting records in two databases

Application


```
write(/db1, "new");  
write(/db2, "new");
```

File System

For each database

- Allocates new data block (bitmap)
- ▶ • Fills data block with user data (data)
- Sets location of data block (inode)

Storage Device



Example: inserting records in two databases

Application


```
write(/db1, "new");  
write(/db2, "new");
```

File System

For each database

- Allocates new data block (bitmap)
- Fills data block with user data (data)
- Sets location of data block (inode)

Storage Device



Example: inserting records in two databases

Application


```
write(/db1, "new");  
write(/db2, "new");
```

File System

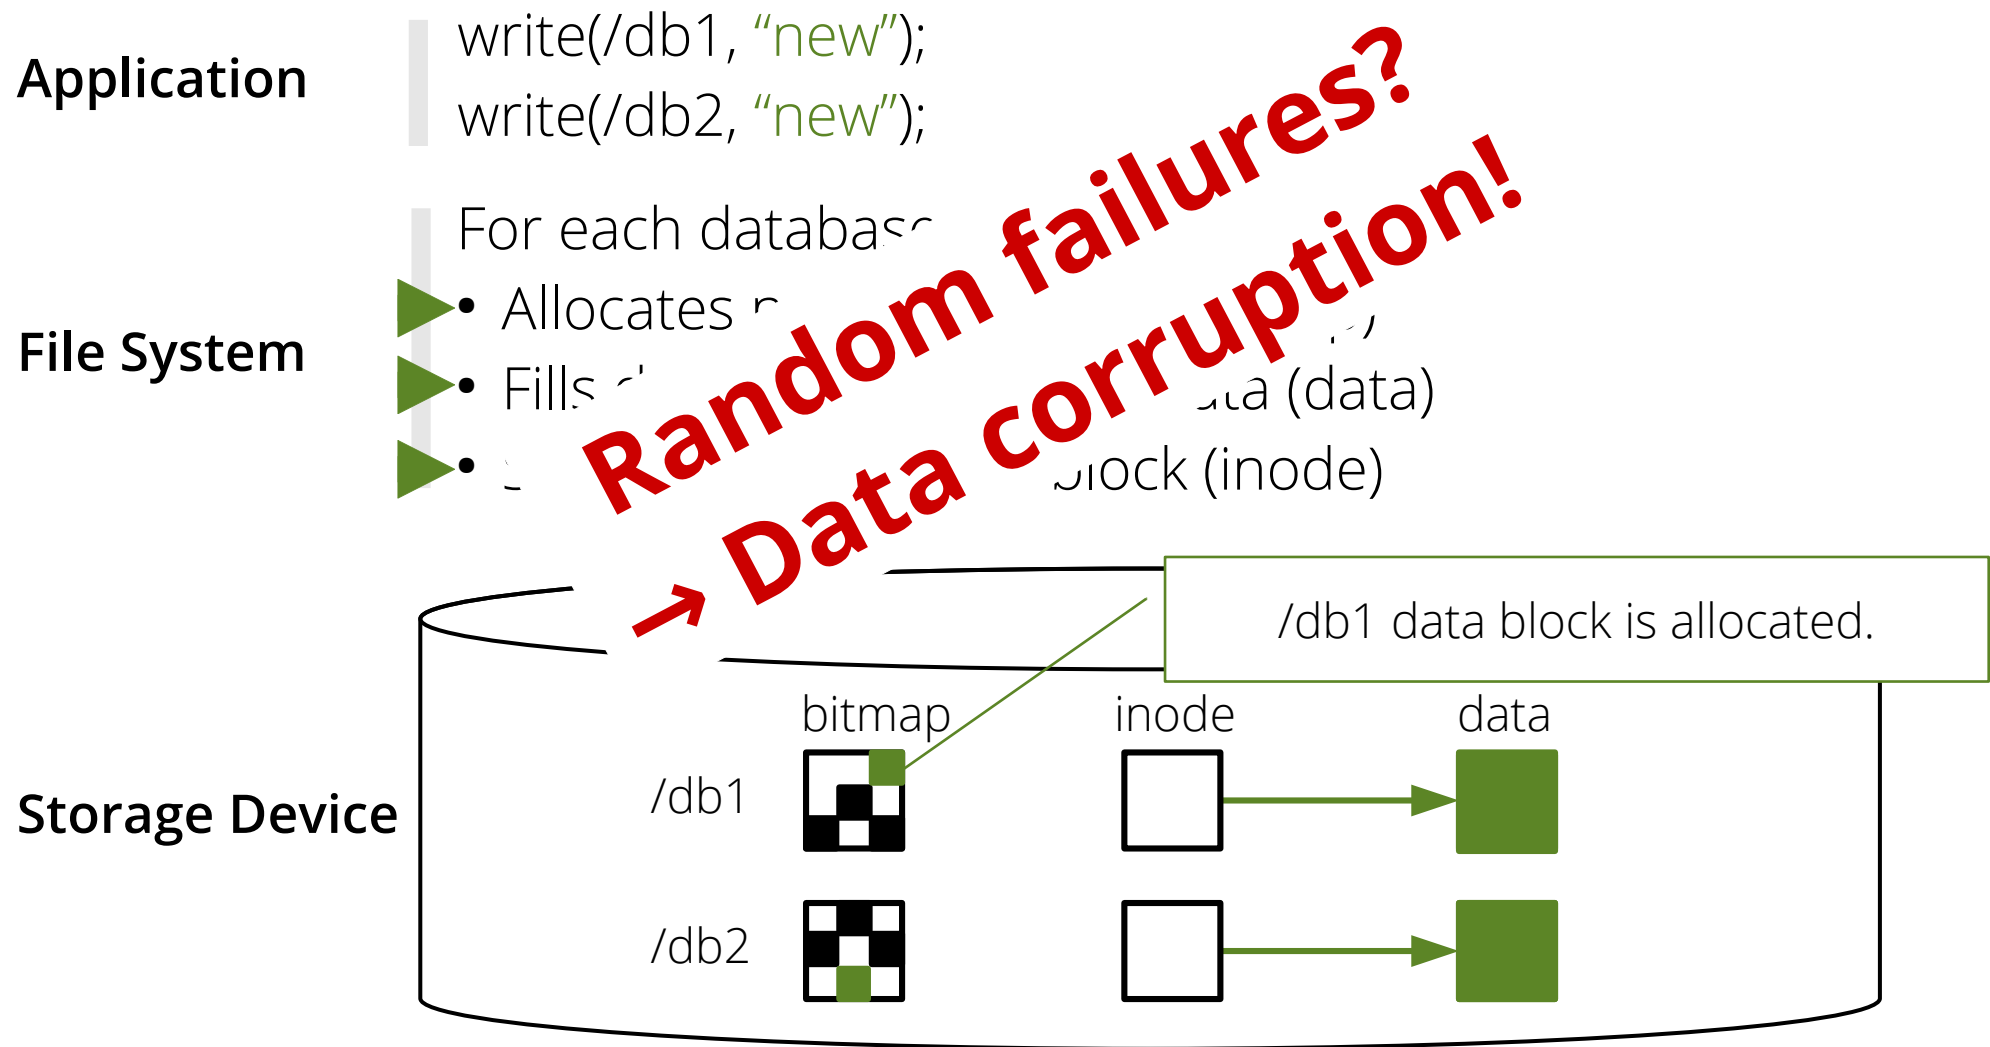
For each database

- Allocates new data block (bitmap)
- Fills data block with user data (data)
- Sets location of data block (inode)

Storage Device



Example: inserting records in two databases




How to Achieve Crash Consistency : The Case of SQLite

```
atomic_update {
```

```
  write(/db1, "new");
```

```
  write(/db2, "new");
```

```
}
```




Logging (i.e., journaling) & Crash Recovery

"Write logs first before writing data in place"

How to Achieve Crash Consistency : The Case of SQLite

```
atomic_update {  
  write(/db1, "new");  
  write(/db2, "new");  
}
```



Logging (i.e., journaling) & Crash Recovery


"Write logs first before writing data in place"

Maintaining three log files

: for each DB and their master
: 3 create() & 3 unlink()

How to Achieve Crash Consistency : The Case of SQLite

```
atomic_update {  
  write(/db1, "new");  
  write(/db2, "new");  
}
```



Logging (i.e., journaling) & Crash Recovery

"Write logs first before writing data in place"

Maintaining three log files


: for each DB and their master
: 3 create() & 3 unlink()

Redundant write

: 7 write()

How to Achieve Crash Consistency : The Case of SQLite

```
atomic_update {  
  write(/db1, "new");  
  write(/db2, "new");  
}
```



Logging (i.e., journaling) & Crash Recovery

"Write logs first before writing data in place"


Ordering & durability
: 11 fsync()

Maintaining three log files
: for each DB and their master
: 3 create() & 3 unlink()

Redundant write
: 7 write()

How to Achieve Crash Consistency : The Case of SQLite

```
atomic_update {  
  write(/db1, "new");  
  write(/db2, "new");  
}
```




```
// create master journal  
open(/master.jnl);  
write(/master.jnl, "/db1,/db2");  
fsync(/master.jnl);  
fsync(/);  
// update db1  
open(/db1.jnl);  
write(/db1.jnl, "old");  
fsync(/db1.jnl);  
fsync(/);  
write(/db1.jnl, "master.jnl");  
fsync(/db1.jnl);  
write(/db1, "new");  
fsync(/db1);
```

```
// update db2  
open(/db2.jnl);  
write(/db2.jnl, "old");  
fsync(/db2.jnl);  
fsync(/);  
write(/db2.jnl, "master.jnl");  
fsync(/db2.jnl);  
write(/db2, "new");  
fsync(/db2);  
// clean up journals  
unlink(/master.jnl);  
fsync(/);  
unlink(/db1.jnl);  
unlink(/db1.jnl);
```

How to Achieve Crash Consistency : The Case of SQLite


```
atomic_update {  
  write(/db1, "new");  
  write(/db2, "new");  
}
```




```
// create master journal  
open(/master.jnl);  
write(/master.jnl, "/db1,/db2");  
fsync(/master.jnl);  
fsync(/);  
// update db1  
open(/db1.jnl);  
write(/db1.jnl, "old");  
fsync(/db1.jnl);  
fsync(/);  
write(/db1.jnl, "master.jnl");  
fsync(/db1.jnl);  
write(/db1, "new");  
fsync(/db1);
```

```
// update db2  
open(/db2.jnl);  
write(/db2.jnl, "old");  
fsync(/db2.jnl);  
fsync(/);  
write(/db2.jnl, "master.jnl");  
fsync(/db2.jnl);  
write(/db2, "new");  
fsync(/db2);  
// clean up journals  
unlink(/master.jnl);  
fsync(/);  
unlink(/db1.jnl);  
unlink(/db1.jnl);
```

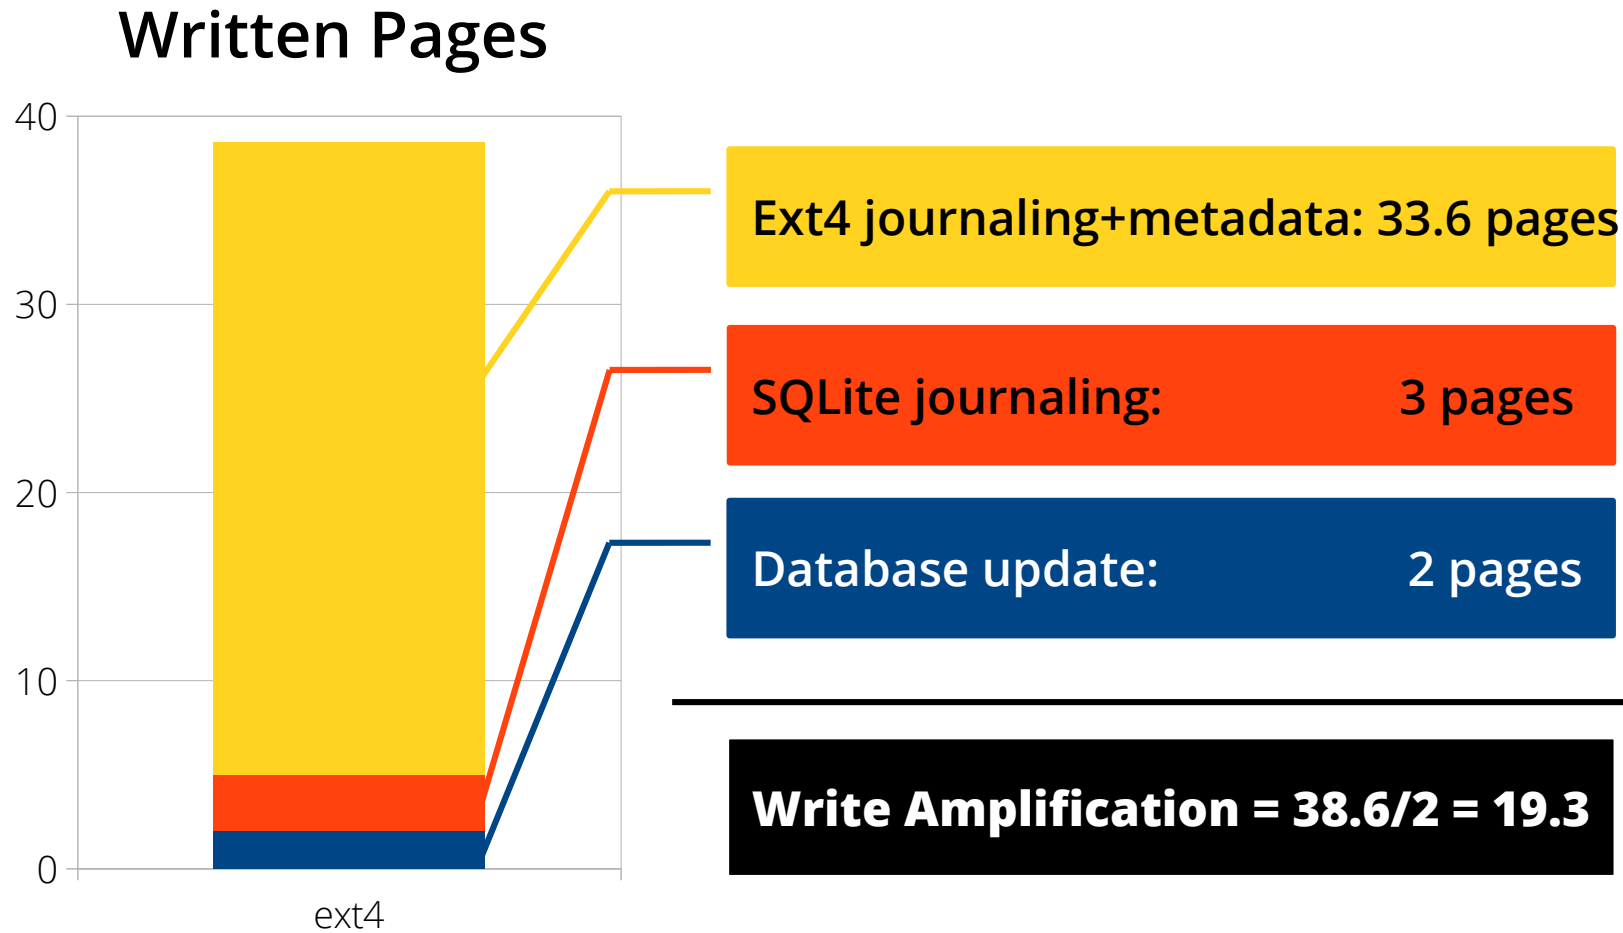

Cost of Crash Consistency : The Case of SQLite



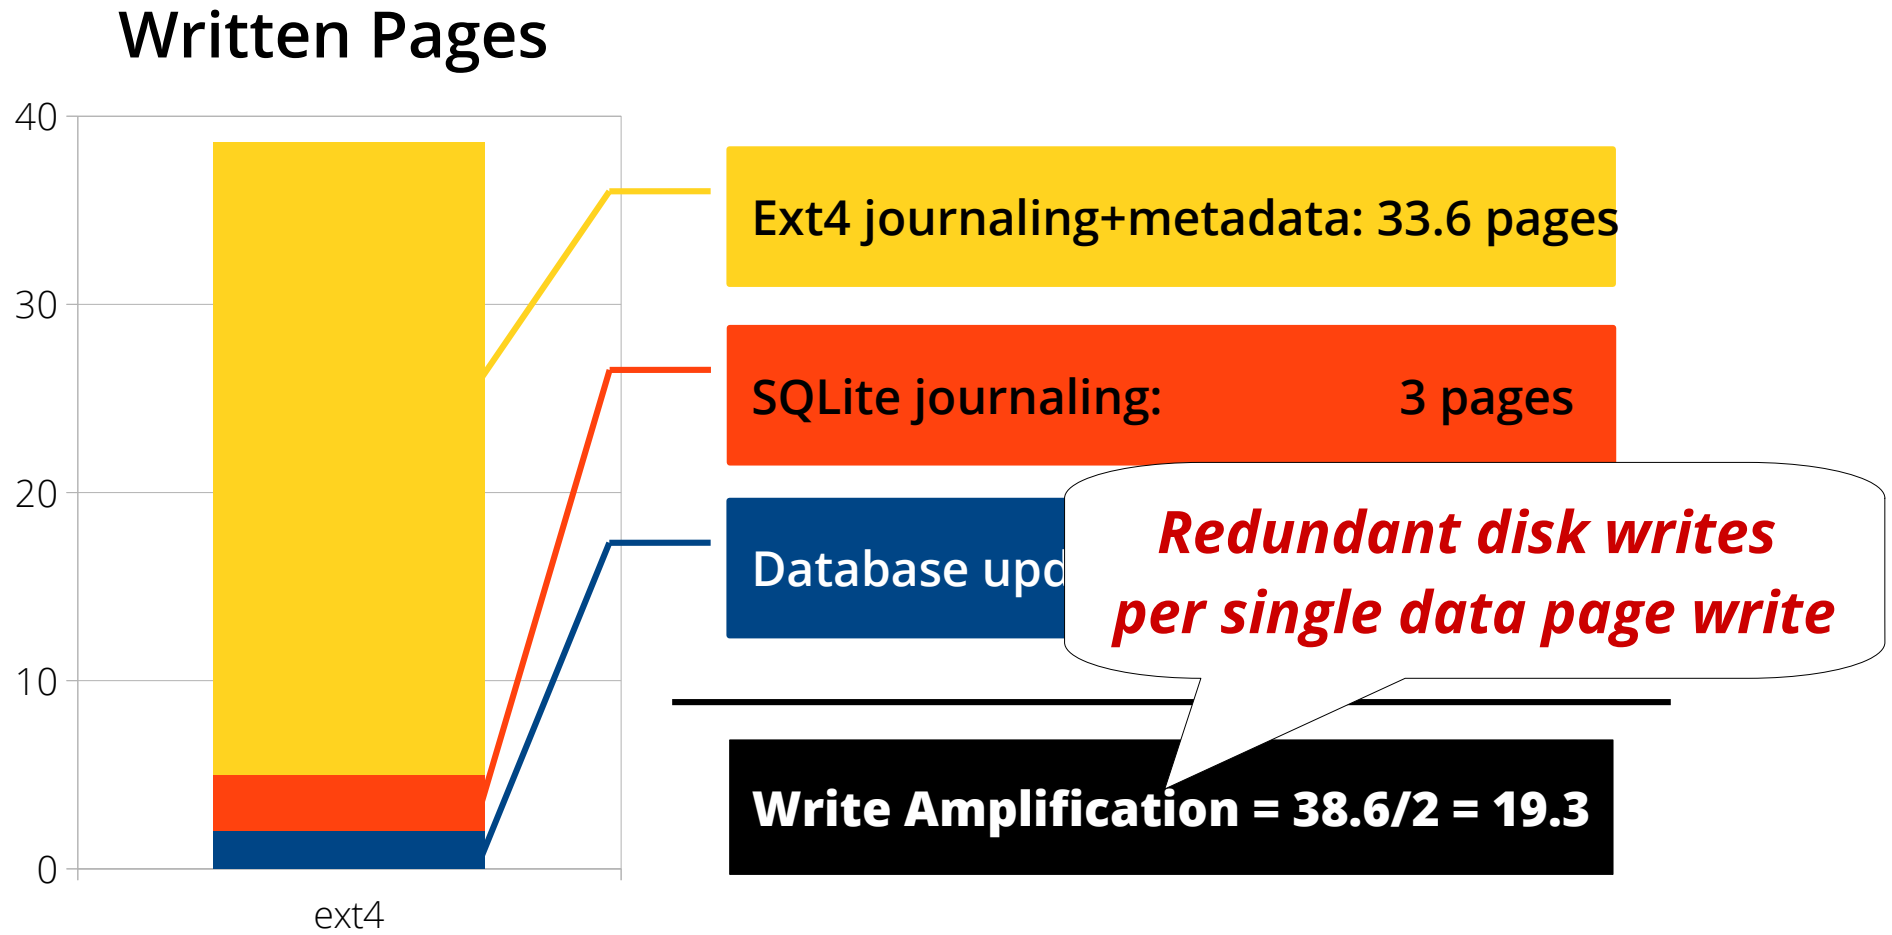
Cost of Crash Consistency : The Case of SQLite




Cost of Crash Consistency : The Case of SQLite



Cost of Crash Consistency : The Case of SQLite



Cost of Crash Consistency : The Case of SQLite



Cost of Crash Consistency : The Case of SQLite

Disk Flush

Notoriously slow!

xsyncfs [Nightingale:OSDI06]

OptFS [Chidambaram:SOSP13]

DuraSSD [Kang:SIGMOD14]

...


10

5

0

ext4

fsync() + FS journaling: 12.7



Problem: complex, redundant
software stack for crash consistency

How can we **simplify** mechanisms for
application crash consistency?

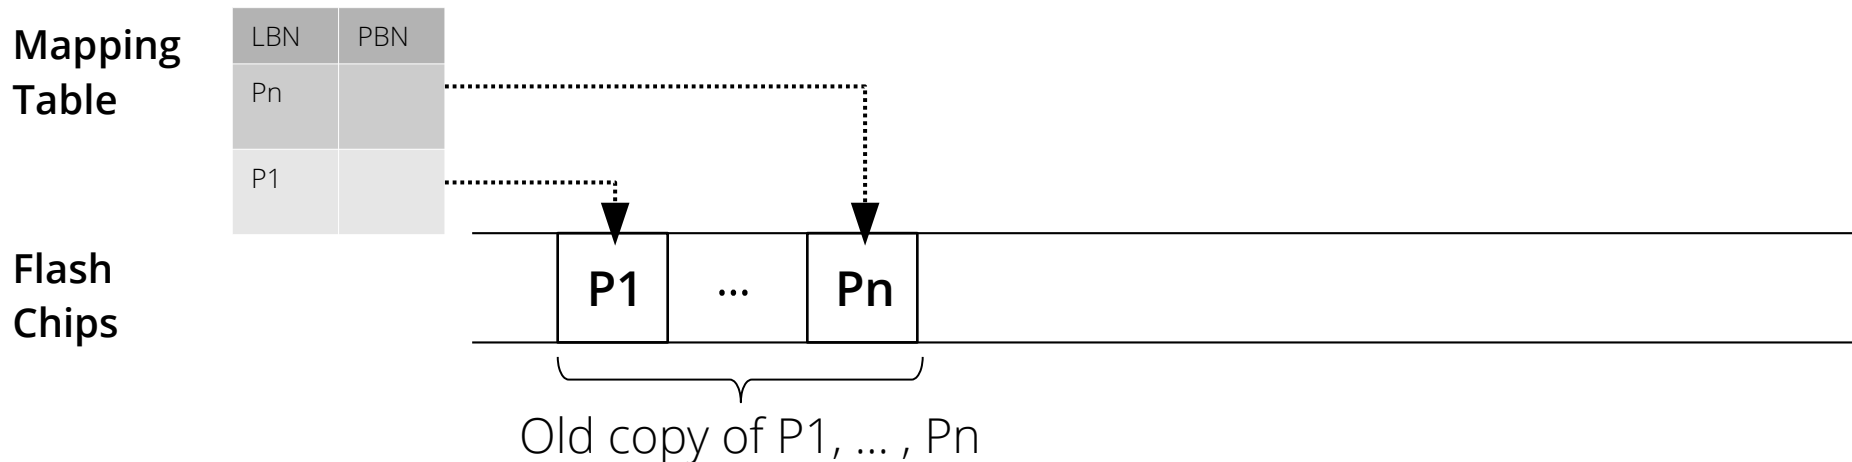
Problem: complex, redundant
software stack for crash consistency

How can we simplify mechanisms for
application crash consistency?

*Can we use **atomic updates of multi pages**
provided by **transactional flash**?*

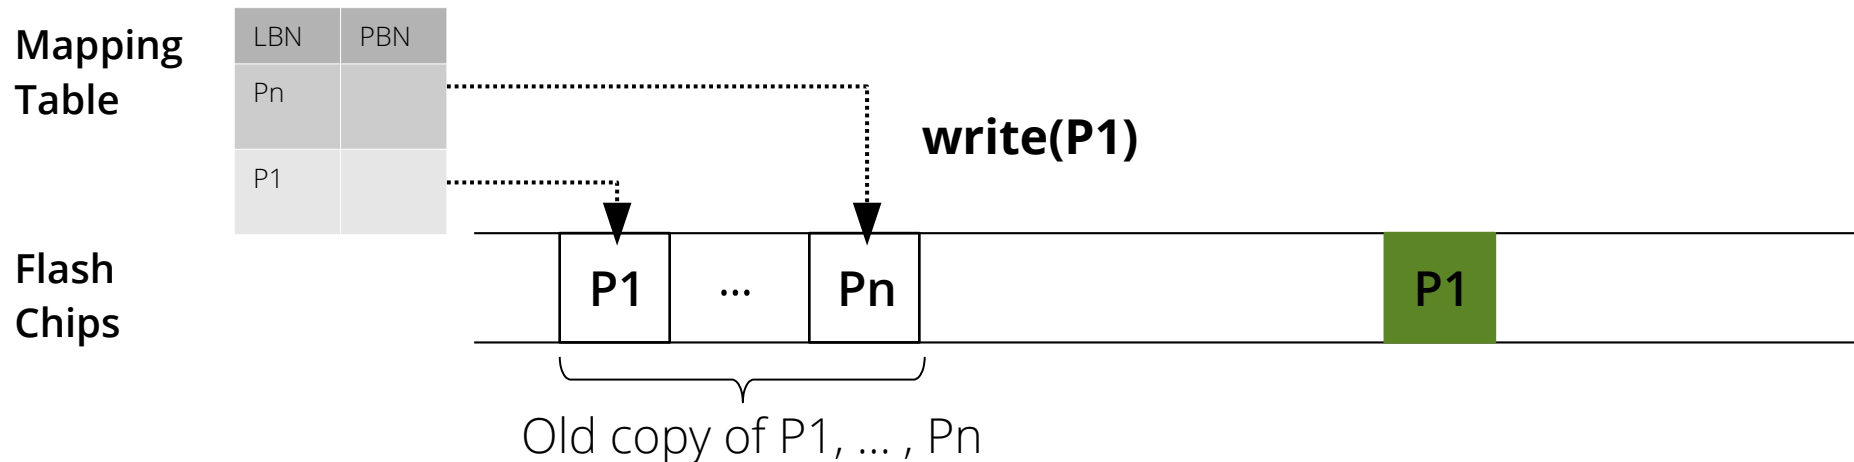
Transactional Flash 101

- NAND Flash SSD
 - No in-place update
 - Log-structured write
 - Mapping table: logical address \rightarrow physical address



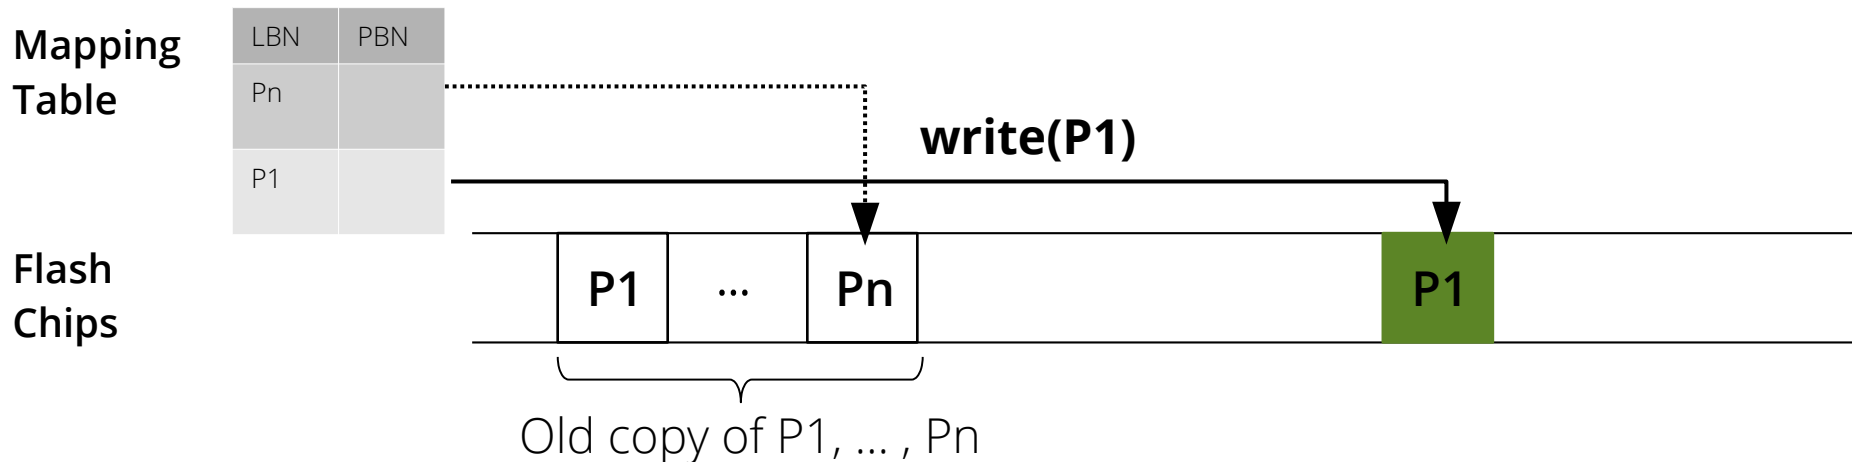
Transactional Flash 101

- NAND Flash SSD
 - No in-place update
 - Log-structured write
 - Mapping table: logical address \rightarrow physical address



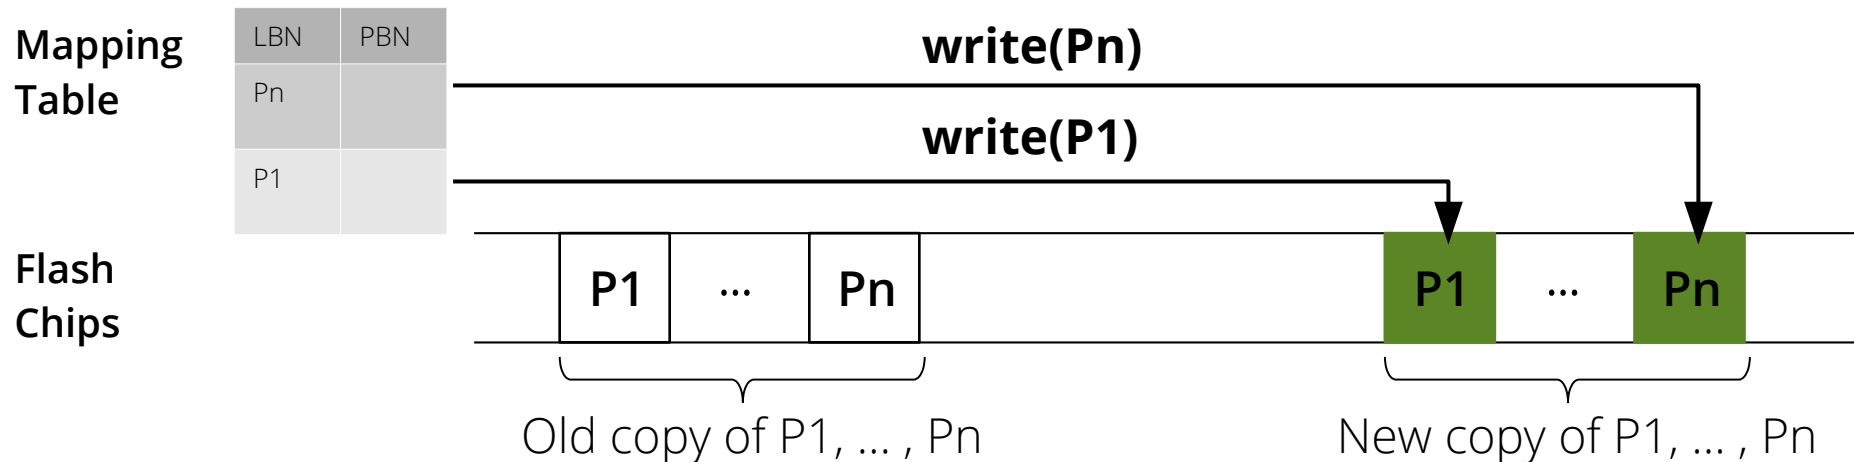
Transactional Flash 101

- NAND Flash SSD
 - No in-place update
 - Log-structured write
 - Mapping table: logical address \rightarrow physical address



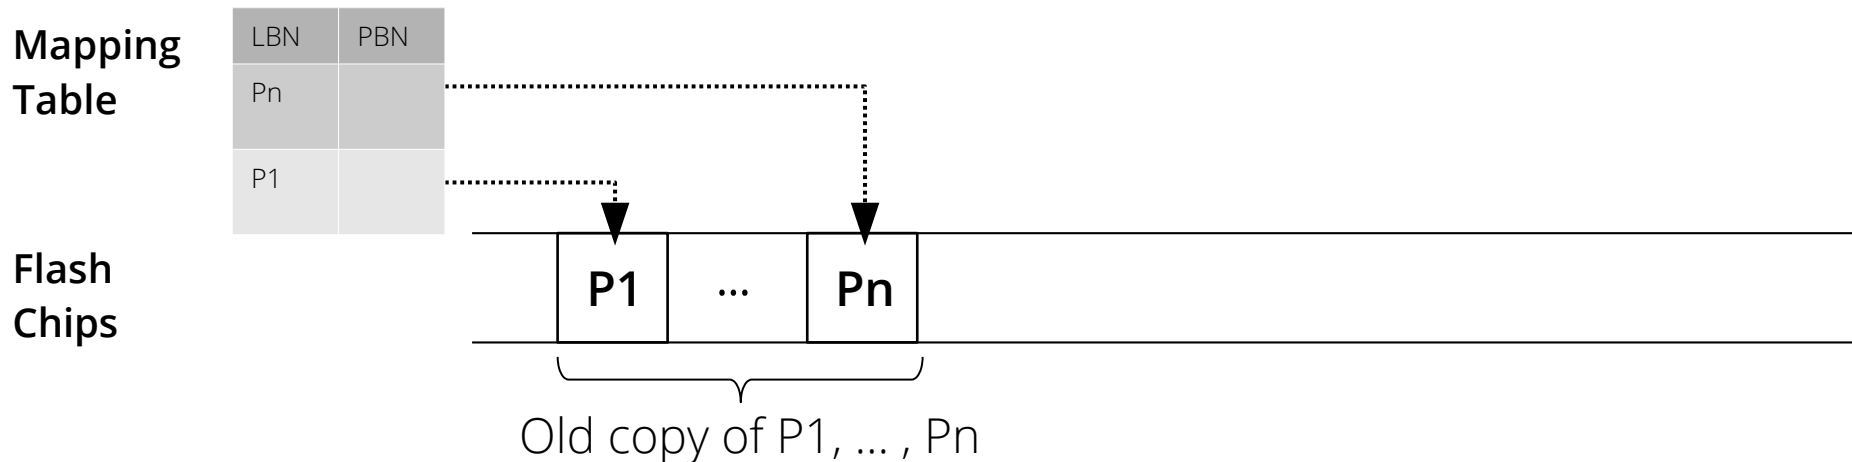
Transactional Flash 101

- NAND Flash SSD
 - No in-place update
 - Log-structured write
 - Mapping table: logical address \rightarrow physical address



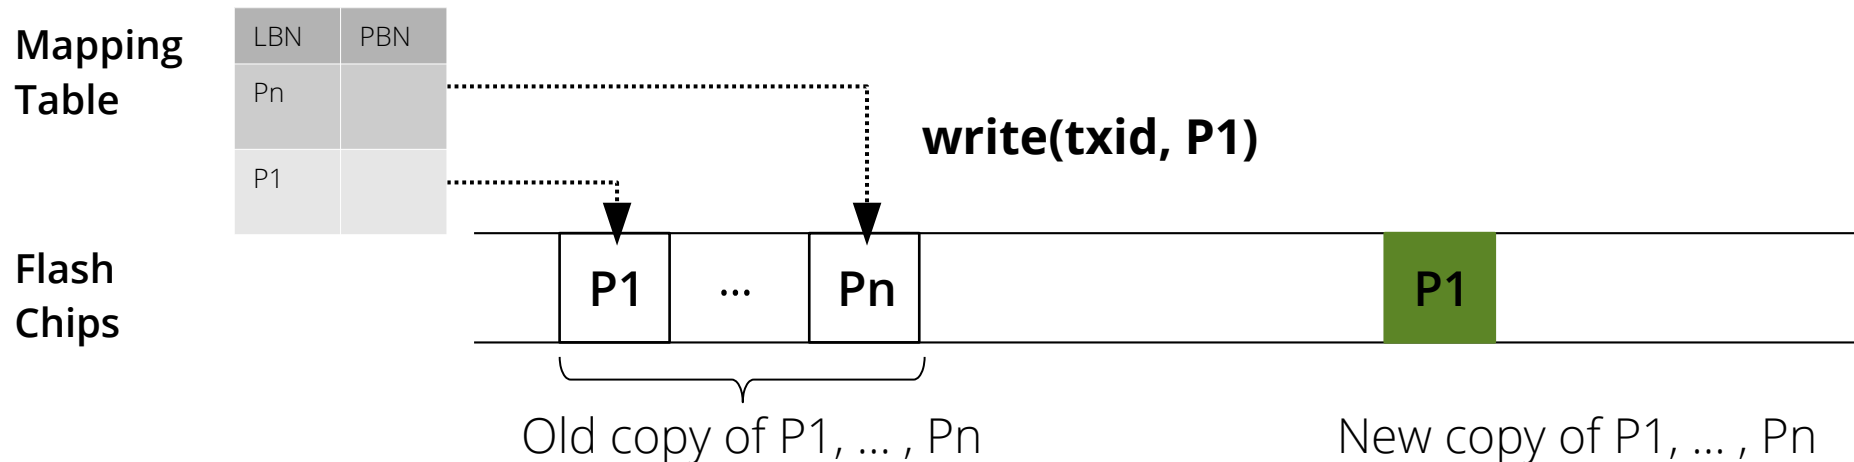
Transactional Flash 101

- Transactional Flash SSD
 - Atomic multi-page write by atomically updating the mapping table at commit request
 - `write(txid, page)`, `commit(txid)`, `abort(txid)`
 - H/W implementation or S/W emulation



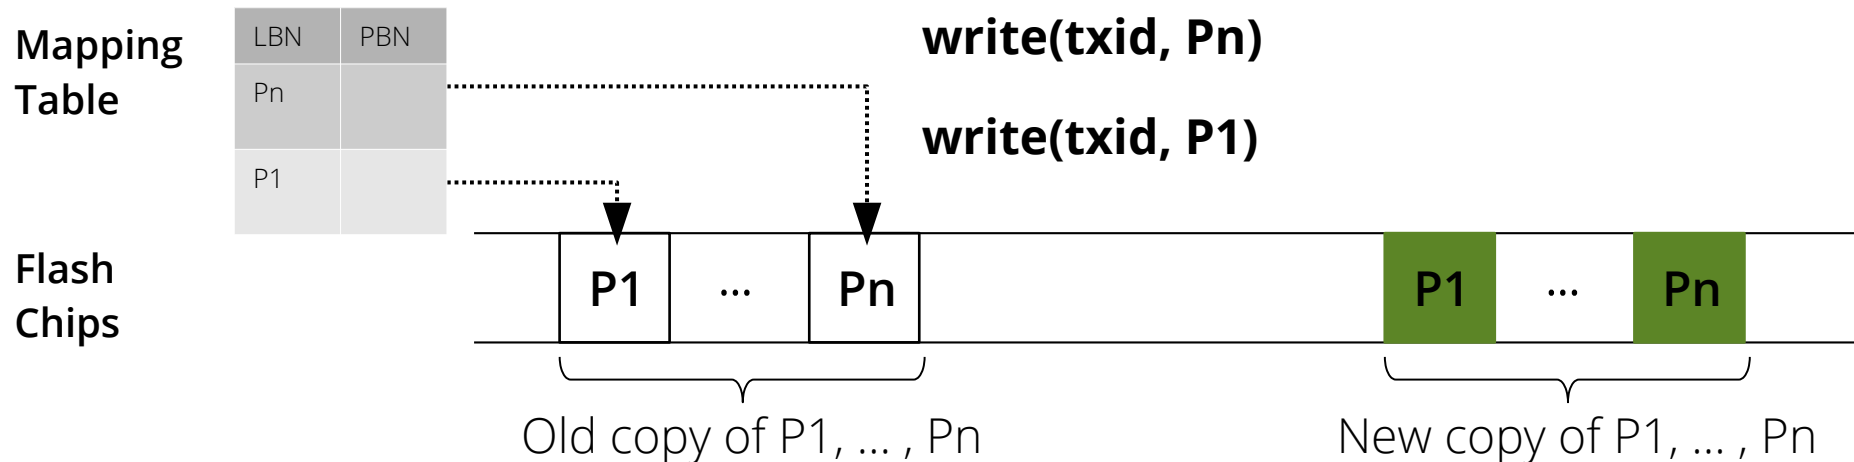
Transactional Flash 101

- Transactional Flash SSD
 - Atomic multi-page write by atomically updating the mapping table at commit request
 - `write(txid, page)`, `commit(txid)`, `abort(txid)`
 - H/W implementation or S/W emulation




Transactional Flash 101

- Transactional Flash SSD
 - Atomic multi-page write by atomically updating the mapping table at commit request
 - `write(txid, page)`, `commit(txid)`, `abort(txid)`
 - H/W implementation or S/W emulation




Transactional Flash 101

- Transactional Flash SSD
 - Atomic multi-page write by atomically updating the mapping table at commit request
 - `write(txid, page)`, `commit(txid)`, `abort(txid)`
 - H/W implementation or S/W emulation



Transactional Flash 101

- Transactional Flash SSD
 - Atomic multi-page write by atomically updating the mapping table at commit request
 - `write(txid, page)`, `commit(txid)`, `abort(txid)`
 - H/W implementation or S/W emulation




Our Solution: CFS, a new file system
using transactional flash

Simplifying
applications' crash consistency
using atomic multi-page write
of transactional flash

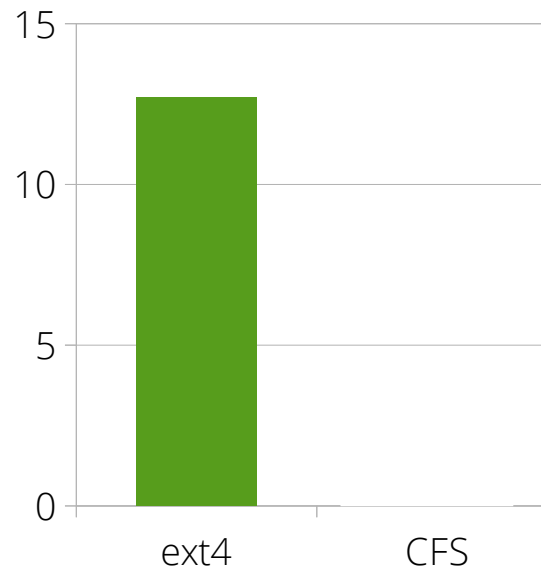
Our Solution: CFS

```
atomic_update {  
  write(/db1, "new");  
  write(/db2, "new");  
}
```

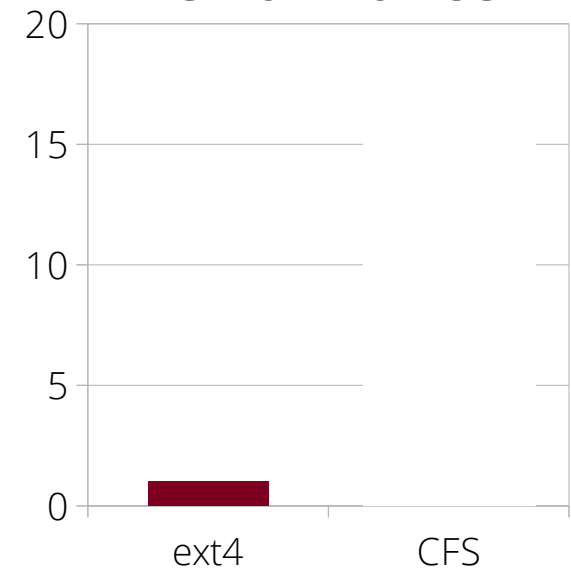
Written Pages



Disk Flush




Performance






Our Solution: CFS

```
atomic_update {  
  write(/db1, "new");  
  write(/db2, "new");  
}
```



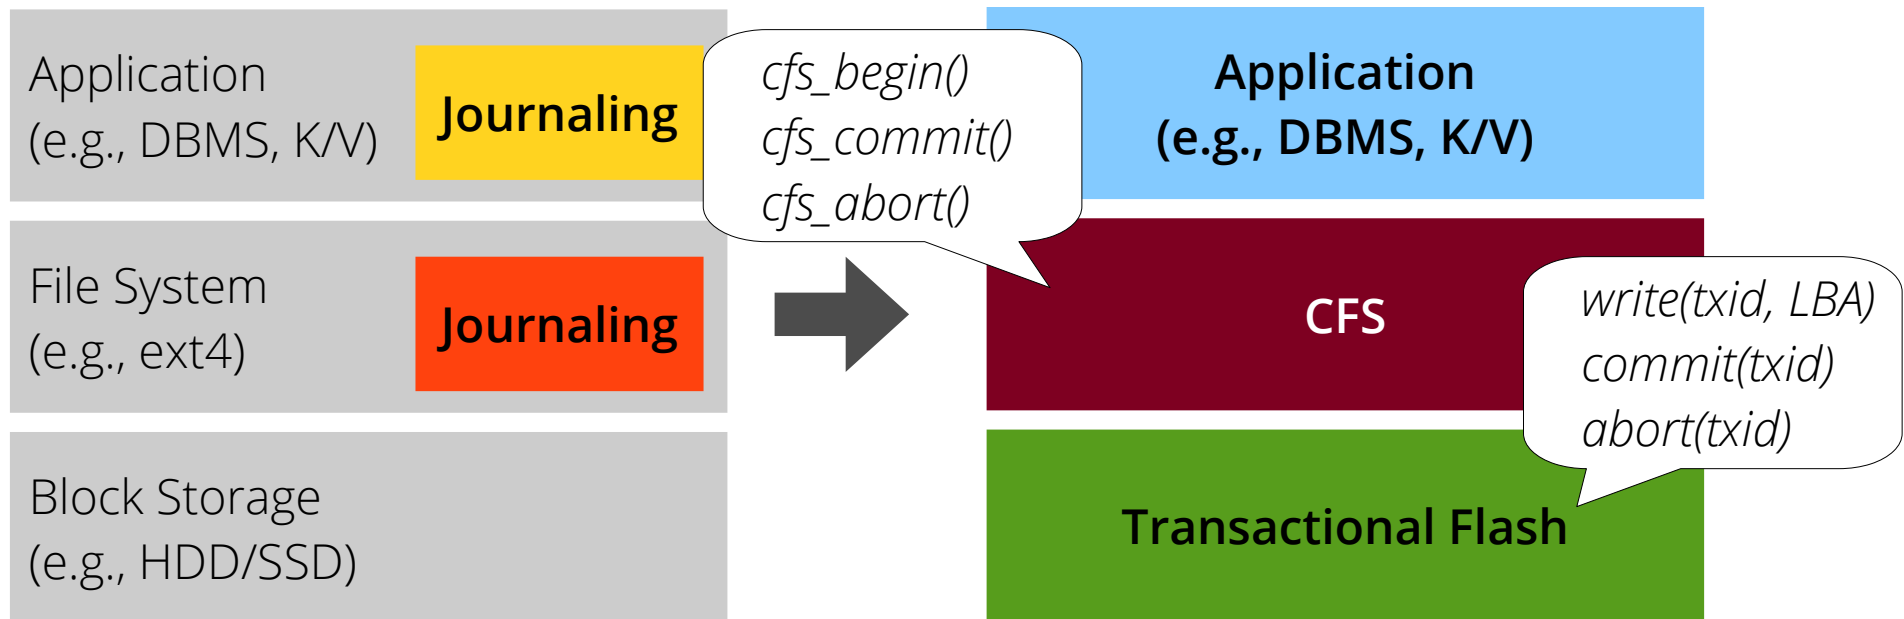
```
+ cfs_begin();  
  write(/db1, "new");  
  write(/db2, "new");  
+ cfs_commit();
```



Outline

- Introduction
- **CFS Design**
- Evaluation
- Conclusion

CFS Architecture



Removing redundant journaling with new primitives provided by transactional flash

Four Challenges

1. Finding a set of pages for atomic updates
2. File system metadata consistency in every case
3. Concurrency control among atomic updates
4. Legacy application support without any modification

#1. Unit of Atomic Write : Atomic Propagation Group

Application

+ *cfs_begin();*

```
write(/db1, "new");
```

```
write(/db2, "new");
```


+ *cfs_commit();*

File System

For each database

- Allocates new data block (bitmap)
- Fills data block with user data (data)
- Sets location of data block (inode)

Transactional Flash



#1. Unit of Atomic Write : Atomic Propagation Group

Application

+ *cfs_begin();*

```
write(/db1, "new");
```

```
write(/db2, "new");
```


+ *cfs_commit();*

File System

For each database

- Allocates new data block (bitmap)
- Fills data block with user data (data)
- Sets location of data block (inode)

Transactional Flash



#1. Unit of Atomic Write : Atomic Propagation Group

Application

+ *cfs_begin();*

```
write(/db1, "new");
```

```
write(/db2, "new");
```


+ *cfs_commit();*

File System

For each database

- Allocates new data block (bitmap)
- Fills data block with user data (data)
- Sets location of data block (inode)

Transactional Flash



#1. Unit of Atomic Write : Atomic Propagation Group

Application

```
+ cfs_begin();
```

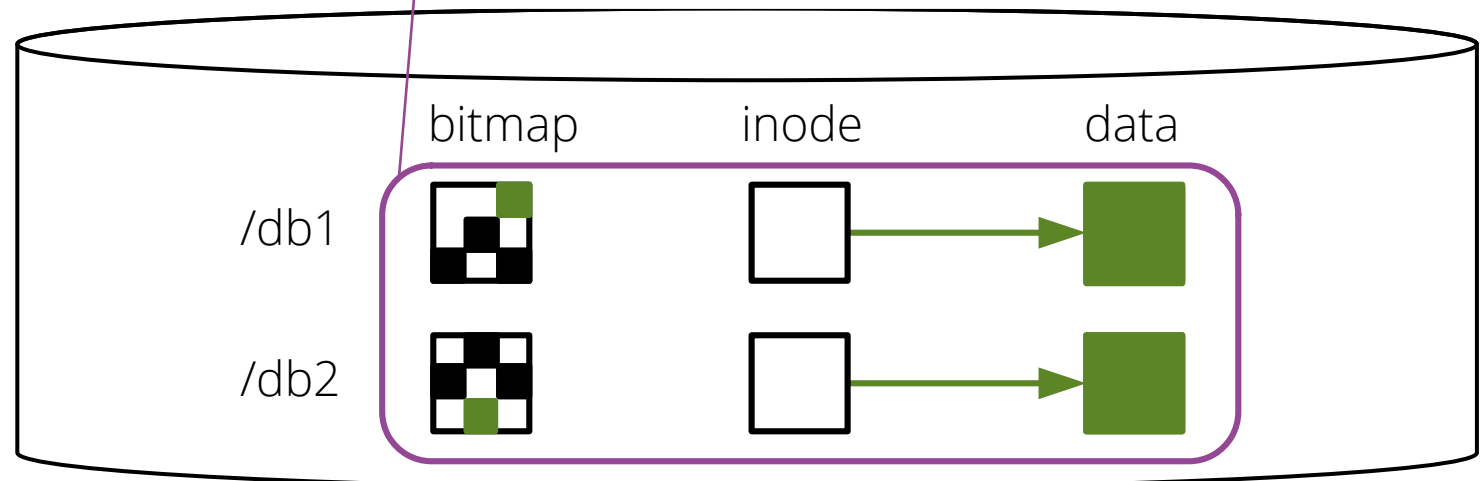
```
write(/db1, "new");
```

```
write(/db2, "new");
```

```
+ cfs_commit();
```

- Updated data and file system metadata pages need to be atomically updated.
- Use atomic multi-page write operations
→ Atomic Propagation Group

Transactional
Flash




Example: Two apps. are running

Application

```
+ cfs_begin();  
  write(/db1, "new");  
  write(/db2, "new");  
+ cfs_commit();
```

Transactional
Flash



Example: Two apps. are running

Application

+ *cfs_begin()*;

write(/db1, "new");

write(/db2, "new");


+ *cfs_commit()*;

+ *cfs_begin()*;

write(/db3, "new");

...

Transactional
Flash



Example: Two apps. are running

Application

+ *cfs_begin()*;

write(/db1, "new");

write(/db2, "new");

+ *cfs_commit()*;

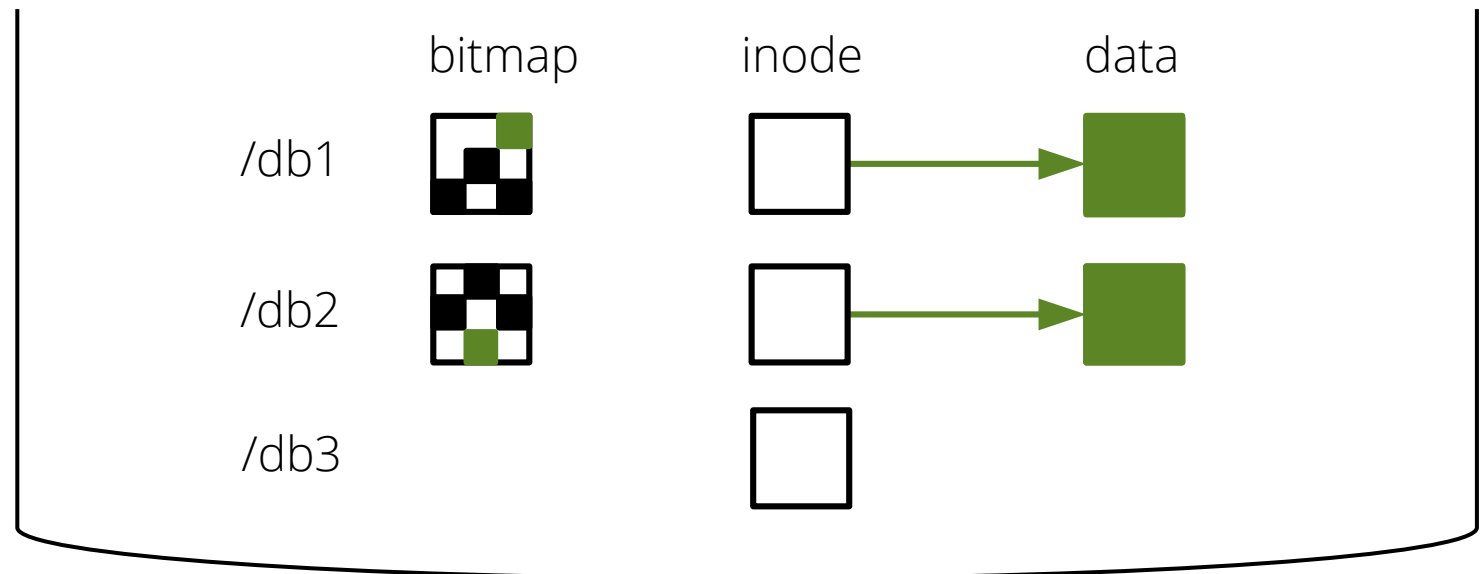
+ *cfs_begin()*;

write(/db3, "new");

...

What if /db3 bitmap locates in the /db2 bitmap?

Transactional
Flash



Example: Two apps. are running

Application

+ *cfs_begin();*

write(/db1, "new");

write(/db2, "new");

+ *cfs_commit();*


+ *cfs_begin();*

write(/db3, "new");

...

What if /db3 bitmap locates in the /db2 bitmap?

Transactional
Flash



Example: Two apps. are running

Application

+ *cfs_begin()*;

write(/db1, "new");

write(/db2, "new");

+ *cfs_commit()*;


+ *cfs_begin()*;

write(/db3, "new");

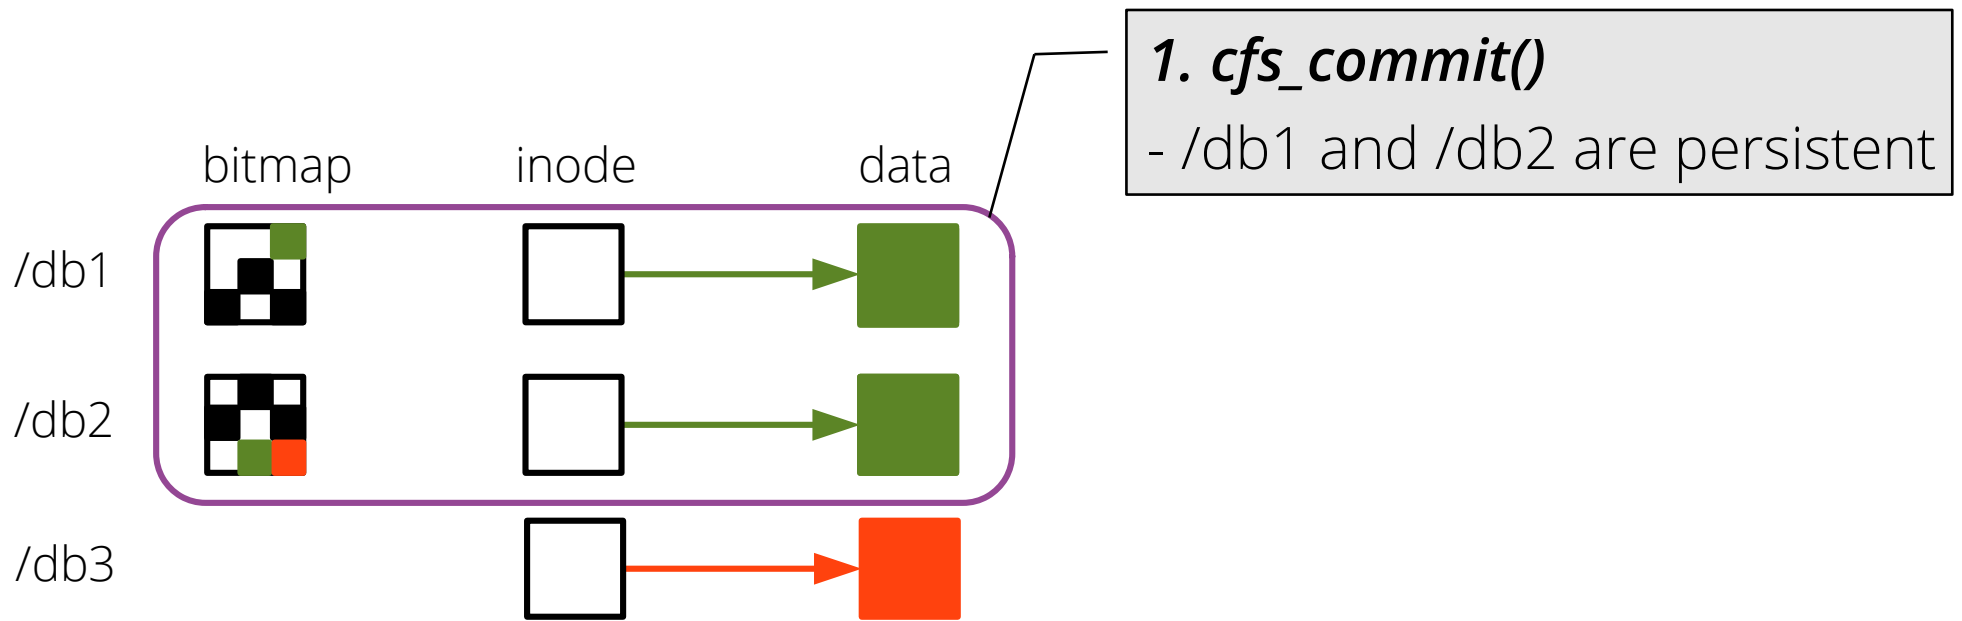
...

What if /db3 bitmap locates in the /db2 bitmap?

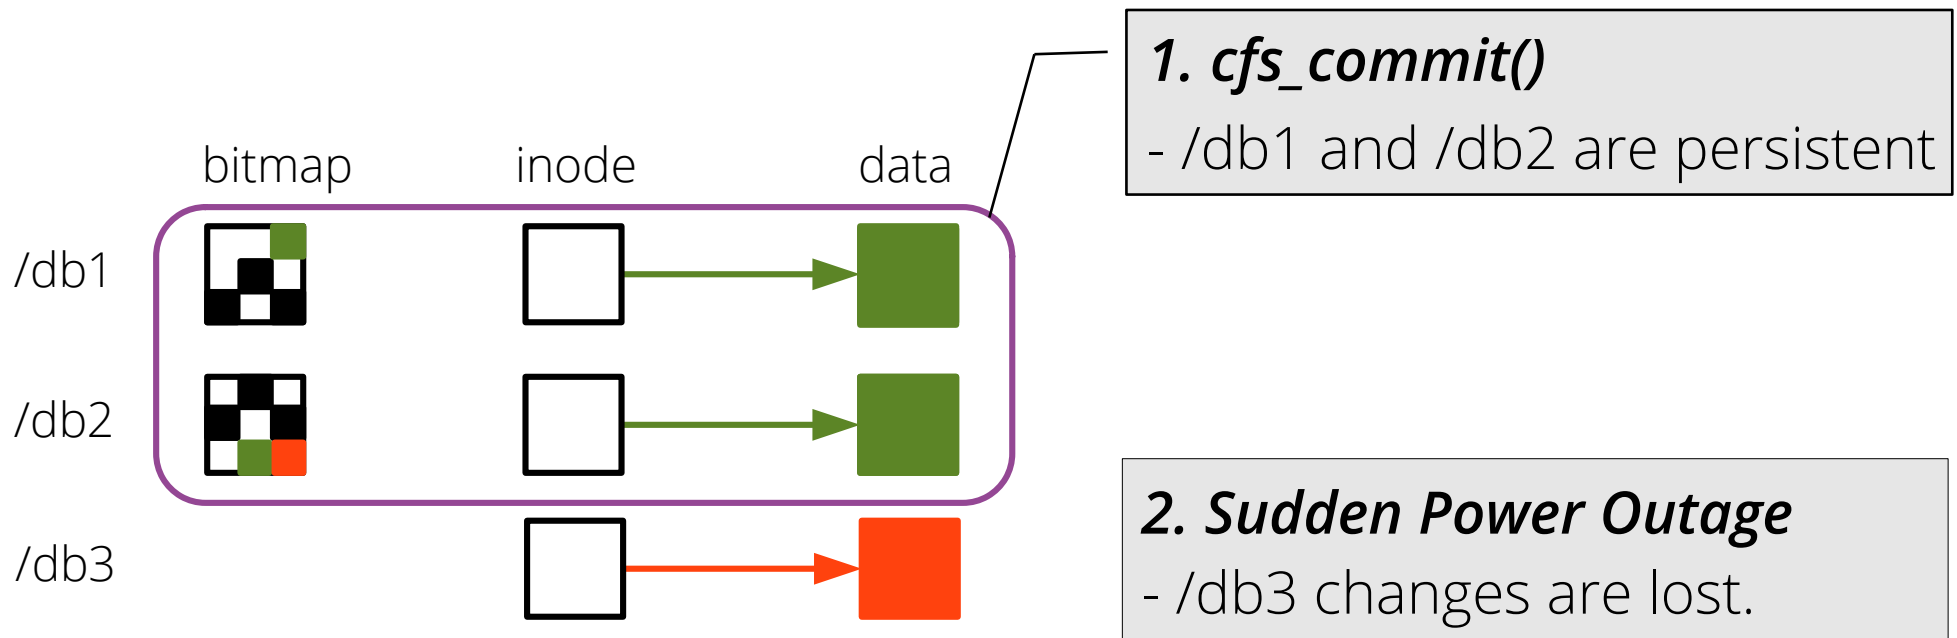
Transactional
Flash



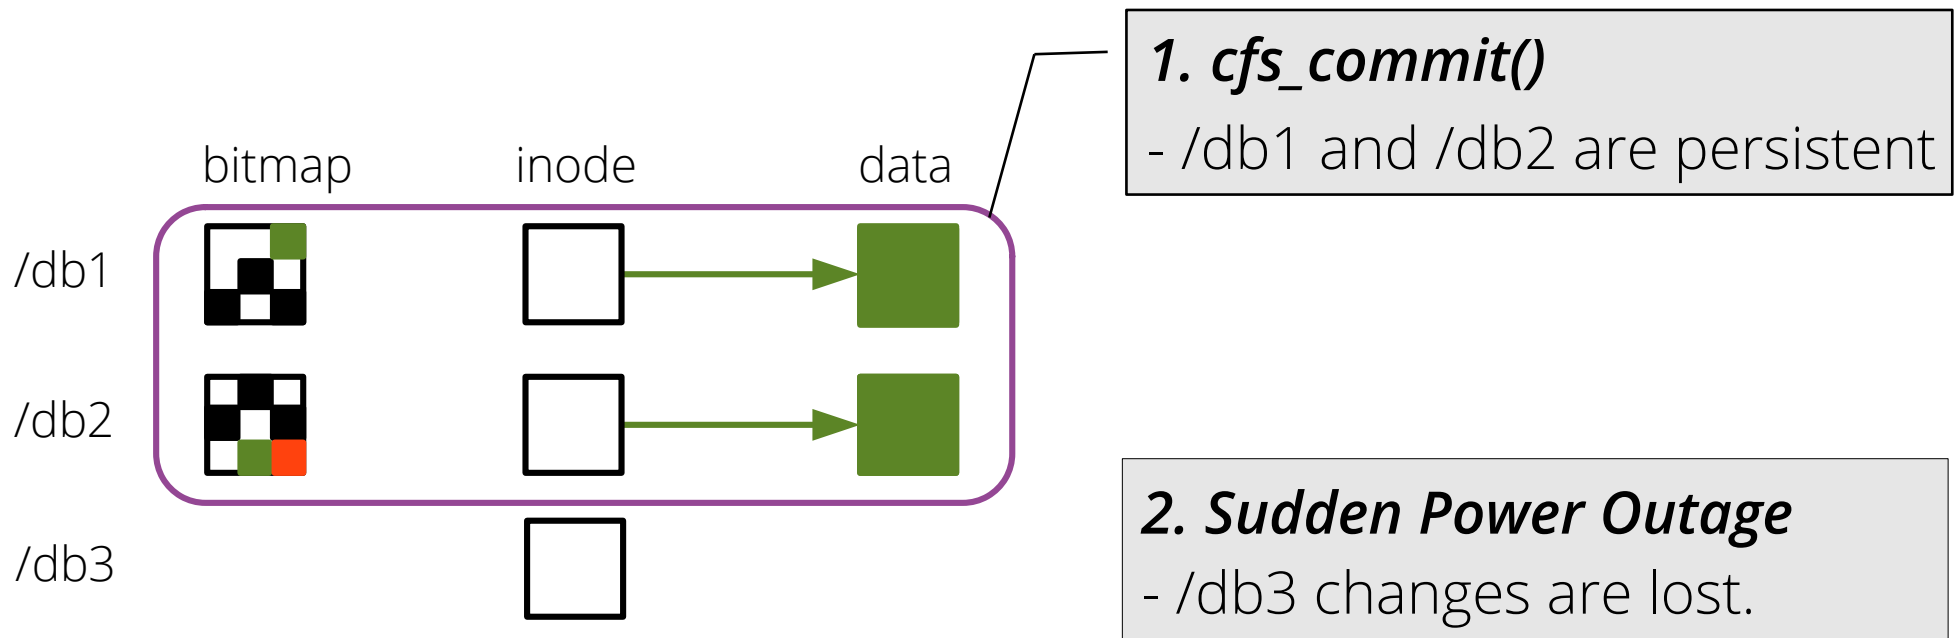
Problem: entangled disk pages



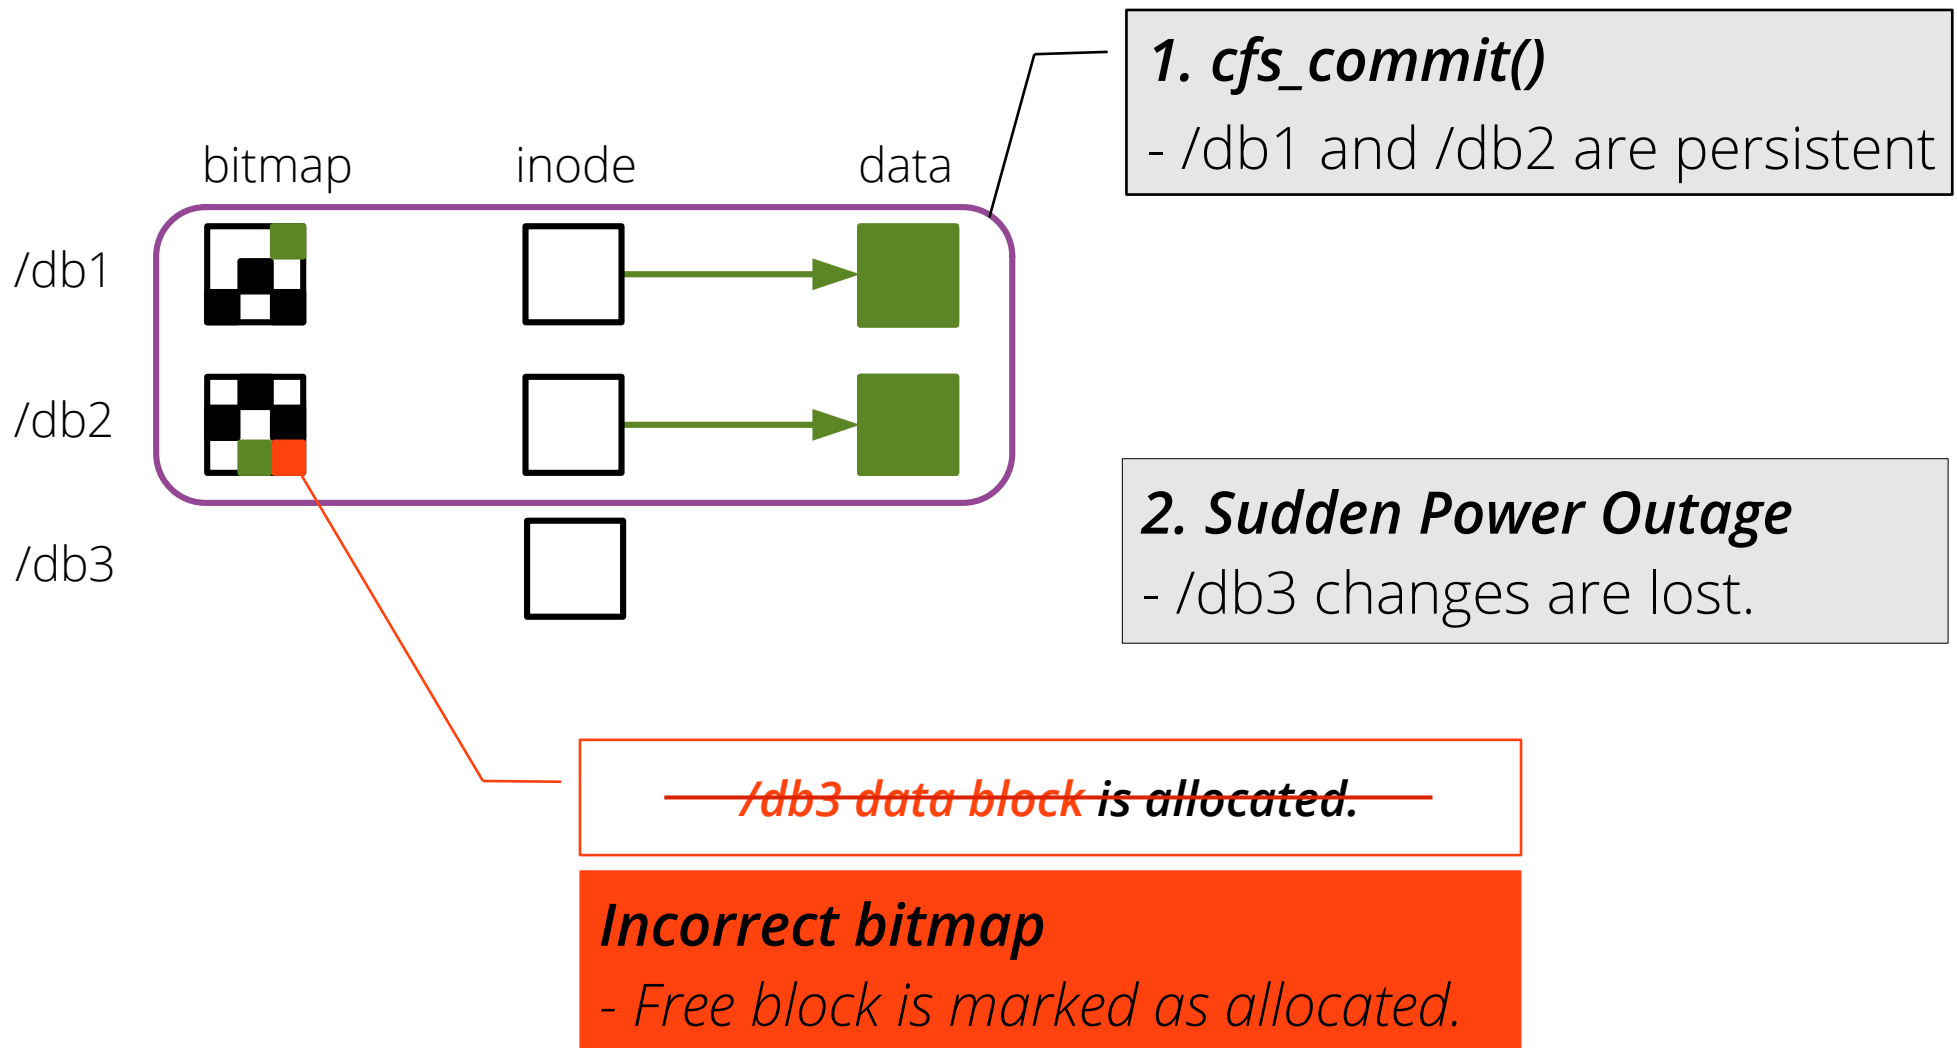
Problem: entangled disk pages



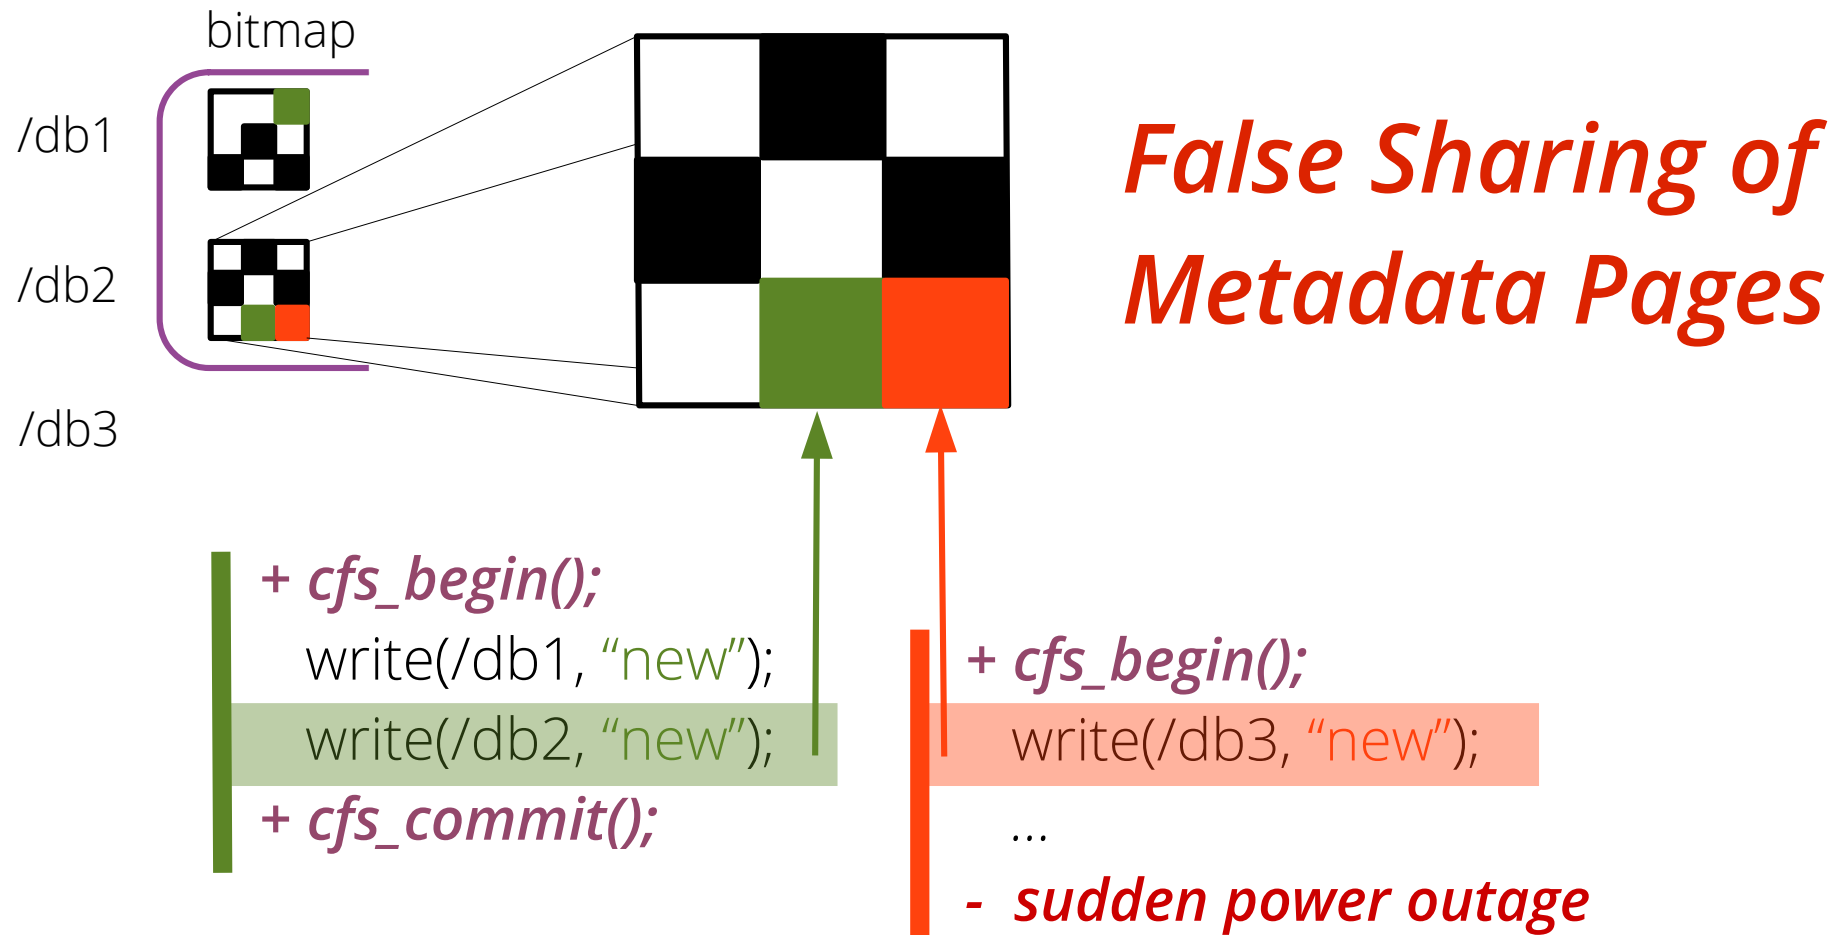
Problem: entangled disk pages



Problem: entangled disk pages



Problem: entangled disk pages



#2. Metadata False Sharing : In-Memory Metadata Logging

- Operational Logging for in-memory metadata change
 - *toggle_bit(free_block_bitmap, LBA)*
 - *sub(free_block_count, 1)*
- Maintain two versions of in-memory metadata to selectively propagate only relevant changes to storage
 - Memory version: on-going modification
 - Storage version: committed version, used for storage IO

REDO or UNDO operational logs

```
cfs_commit() {  
    storage version += REDO(logs);  
    write(txid, storage version);  
    commit(txid);  
}
```

```
cfs_abort() {  
    memory version -= UNDO(logs);  
    abort(txid);  
}
```

Example: Two apps. are running


**In-Memory
Metadata**

Memory
Version



bitmap




Storage
Version



**Operational Log
for Atomic Propagation Group**

App 1	App 2
Turn on a bit 	Turn on a bit 

**Transactional
Flash**



Example: Two apps. are running


**In-Memory
Metadata**

Memory
Version



bitmap




Storage
Version



**Operational Log
for Atomic Propagation Group**

App 1	App 2
Turn on a bit 	Turn on a bit 
<code>cfs_commit()</code>	

**Transactional
Flash**



Example: Two apps. are running


In-Memory Metadata

Memory Version



bitmap




Storage Version



Operational Log for Atomic Propagation Group

App 1	App 2
<input checked="" type="checkbox"/> Turn on a bit 	Turn on a bit 
<code>cfs_commit()</code>	

Transactional Flash




Example: Two apps. are running

**In-Memory
Metadata**

Memory
Version





Storage
Version



Bitmap for db3 is not written.

**Operational Log
for Atomic Propagation Group**

App 1	App 2
✓ Turn on a bit 	Turn on a bit 
✓ <code>cfs_commit()</code>	

**Transactional
Flash**




Example: Two apps. are running

In-Memory Metadata

Memory Version





Storage Version



Bitmap for db3 is not written.

Operational Log for Atomic Propagation Group

App 1	App 2
✓ Turn on a bit 	Turn on a bit 
✓ cfs_commit()	cfs_abort()

Transactional Flash



Example: Two apps. are running


In-Memory
Metadata

Memory
Version

bitmap





Storage
Version



Bitmap for db3 is reverted.

Operational Log
for Atomic Propagation Group

App 1	App 2
✓ Turn on a bit 	⊖ Turn on a bit 
✓ cfs_commit()	⊖ cfs_abort()

Transactional
Flash




Example: Two apps. are running

In-Memory
Metadata

Memory
Version



Storage
Version



Bitmap for db3 is reverted.

Operational Log
for Atomic Propagation Group

App 1	App 2
Turn on a bit	Turn on a bit
<code>cfs_commit()</code>	<code>cfs_abort()</code>

Transactional
Flash




*Selective Propagation
of Metadata Change*

#3. Isolation? Concurrency Control?

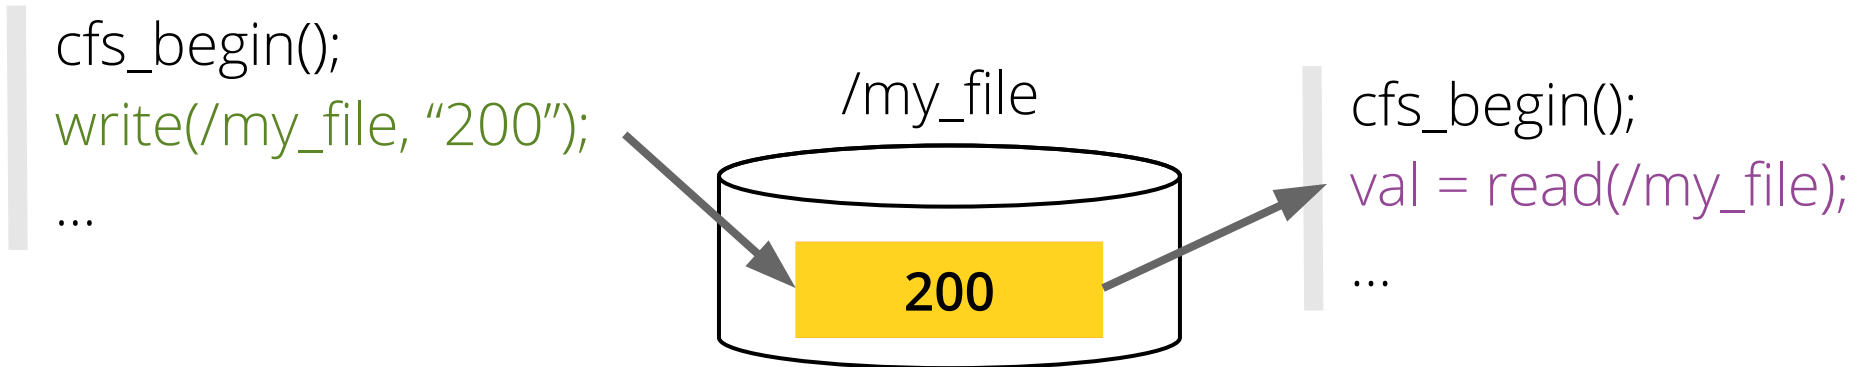
: Leave It to Application

```
cfs_begin();  
write(/my_file, "200");  
...
```



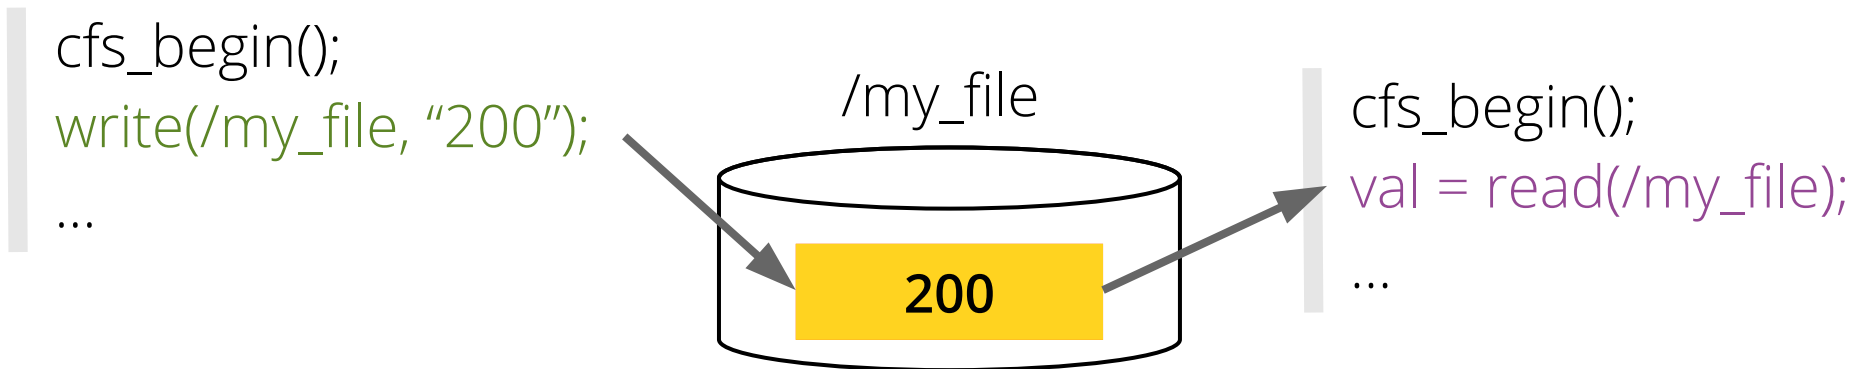
#3. Isolation? Concurrency Control?

: Leave It to Application



#3. Isolation? Concurrency Control?

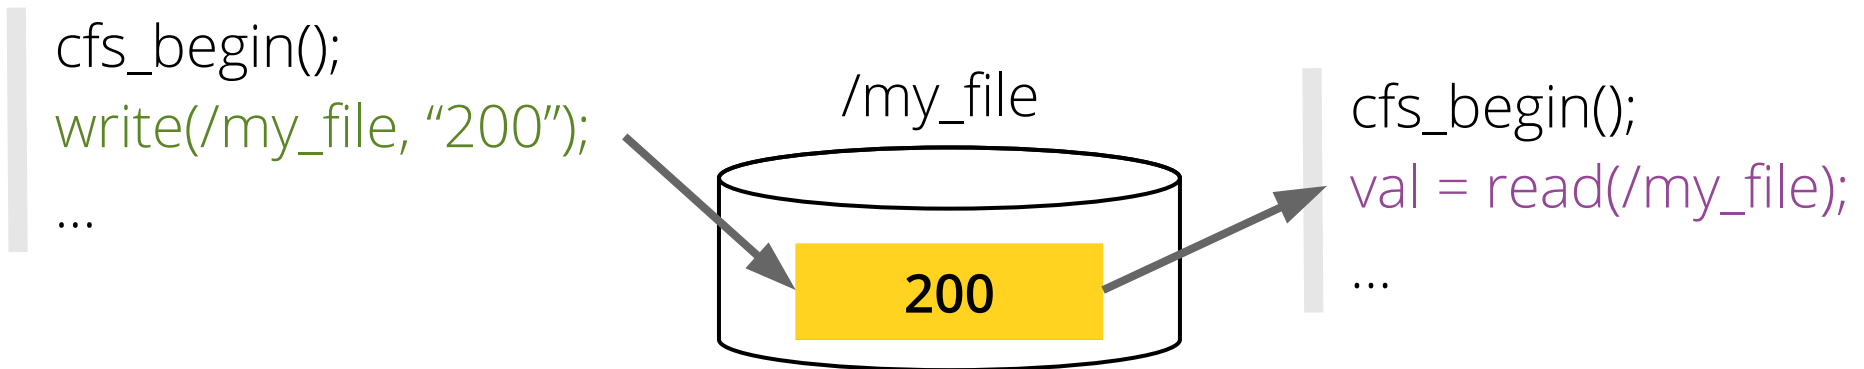
: Leave It to Application



What *val* should be either of **100** or **200**?

#3. Isolation? Concurrency Control?

: Leave It to Application



What *val* should be either of **100** or **200**?

It depends on application semantics.


#3. Isolation? Concurrency Control?

: Leave It to Application



```
> START TRANSACTION  
> UPDATE account  
  SET deposit='200'  
> ...  
> COMMIT
```

/bank_account



```
> START TRANSACTION  
> SELECT deposit  
  FROM account  
> ...  
> COMMIT
```

Deposit of bank account must be 200.


#3. Isolation? Concurrency Control?

: Leave It to Application



```
> START TRANSACTION
> UPDATE account
  SET deposit='200'
> ...
> COMMIT
```

/bank_account




```
> START TRANSACTION
> SELECT deposit
  FROM account
> ...
> COMMIT
```

Deposit of bank account must be 200.



```
KV→begin_transaction(...);
KV→set("account.likes",
      "200");
...
KV→end_transaction(...);
```

/sns_account



```
KV→begin_transaction(...);
KV→get("account.deposit");
...
...
KV→end_transaction(...);
```

Likes of SNS account can be either of 100 or 200.

#3. Isolation? Concurrency Control?

: Leave It to Application

- Isolation and crash-consistency are orthogonal.
 - Even the SQL standard defines four different isolation levels.
- CFS does not provide its own concurrency control mechanism.
 - If needed, use existing synchronization primitives (e.g., mutex, RW lock, etc).

#4. Legacy Application Support

- System-Wide Atomic Propagation Group
 - Every update from legacy applications belongs.
 - Automatically committed by sync() or page flusher


***CFS supports legacy applications
without any modification.***

Outline

- Introduction
- CFS Design
- **Evaluation**
- Conclusion

Implementation

- CFS
 - 5.8k LoC modification of ext4 on Linux 3.10
 - Capture logs by inserting 182 places in ext4
- Transactional Flash
 - OpenSSD: 8KB, 128 pages/block, 8GB w/ SATA2
 - X-FTL/SSD [Kang:SIGMOD'13]



Real Application & Workloads

Mobile
Database



+ RLBench
+ Facebook

Text
Editor



SQL
Database



+ SysBench
+ LinkBench

Key/Value
Store



+ kctreetest
+ db_bench

Package
Installer



Real Application & Workloads

Mobile
Database



+ RLBench
+ Facebook

Text
Editor



CFS is easy-to-use!

: 317 LoC changes out of 3.5 million

Key/Value
Store



+ kctreetest
+ db_bench


Package
Installer




SQLite

+ Facebook App. SQL Trace


Write Amplification



Disk Flush




Performance




Ext4: ordered journal mode
SQLite: rollback journal mode

MariaDB + LinkBench


Write Amplification



Disk Flush




Performance




Ext4: ordered journal mode

KyotoCabinet + db_bench


Write Amplification



Disk Flush



Performance



Ext4: ordered journal mode

Conclusion

- Current mechanisms for crash consistency is complex, slow, and error-prone.
- CFS simplifies application's crash consistency using transactional flash.
 - Atomic propagation group
 - In-memory metadata logging
- Our evaluation shows
 - Less write: 2 ~ 4x ↓
 - Less disk flush: 3 ~ 17x ↓
 - Higher performance: 2 ~ 5x ↑

Thank you!

Changwoo Min

changwoo@gatech.edu

Woon-Hak Kang[†], Taesoo Kim, Sang-Won Lee[†], Young Ik Eom[†]



Georgia Institute of Technology

[†]Sungkyunkwan University



Questions?