



# **Practical Exploit Mitigation Against Code Re-Use Attacks and System Call Abuse**

**Christopher Jelesnianski**

**April 13, 2022**

**Committee:**

Changwoo Min (Chair)

Yeongjin Jang, Danfeng Yao, Wenjie Xiong, & Haibo Zeng



# Outline

- Motivation
- Background
- Contributions
- Future Work
- Summary

# Outline

- **Motivation**
- Background
- Contributions
- Future Work
- Summary

# Thesis Statement

Exploit mitigation techniques need to **shift focus** to not only provide an **effective solution** to the latest threat, but also to be a **practical defense** that can be **readily deployed**


# Mainstream Adoption of Mitigation Techniques

- PaX Linux Address Space Layout Randomization (**ASLR**) (2001)
- Linux **Kernel Stack ASLR** (2002)
- No-Execute Bit (**NX-bit**)
  - Linux 2.6.8 (2004)
  - Windows XP SP2 (2004) -- Data Execution Prevention (**DEP**)
- System Call Whitelisting (**seccomp**)
  - Linux 2.6.12 - (2005)
  - Linux 2.6.23 - (2007)
  - seccomp mode 2 - Linux 3.5 (2012)
  - seccomp eBPF - Linux 3.8 (2013)
- GCC & Clang **Control-Flow Integrity & SafeStack** (USENIX Security 2014)

That's about it...

# Challenges in Exploit Mitigation Design

Why is Practical Exploit Mitigation Design Difficult to Achieve?



## Non-Negligible Performance

- **Complex** or **frequent** verification is hard to make **fast**



## Attack Complexity

- Modern attacks require **more code components** to be **guarded**




## Limited Scalability

- Designing mechanisms with **negligible system impact** is hard



## Fragility

- **Perfect analysis** is challenging to achieve in practice

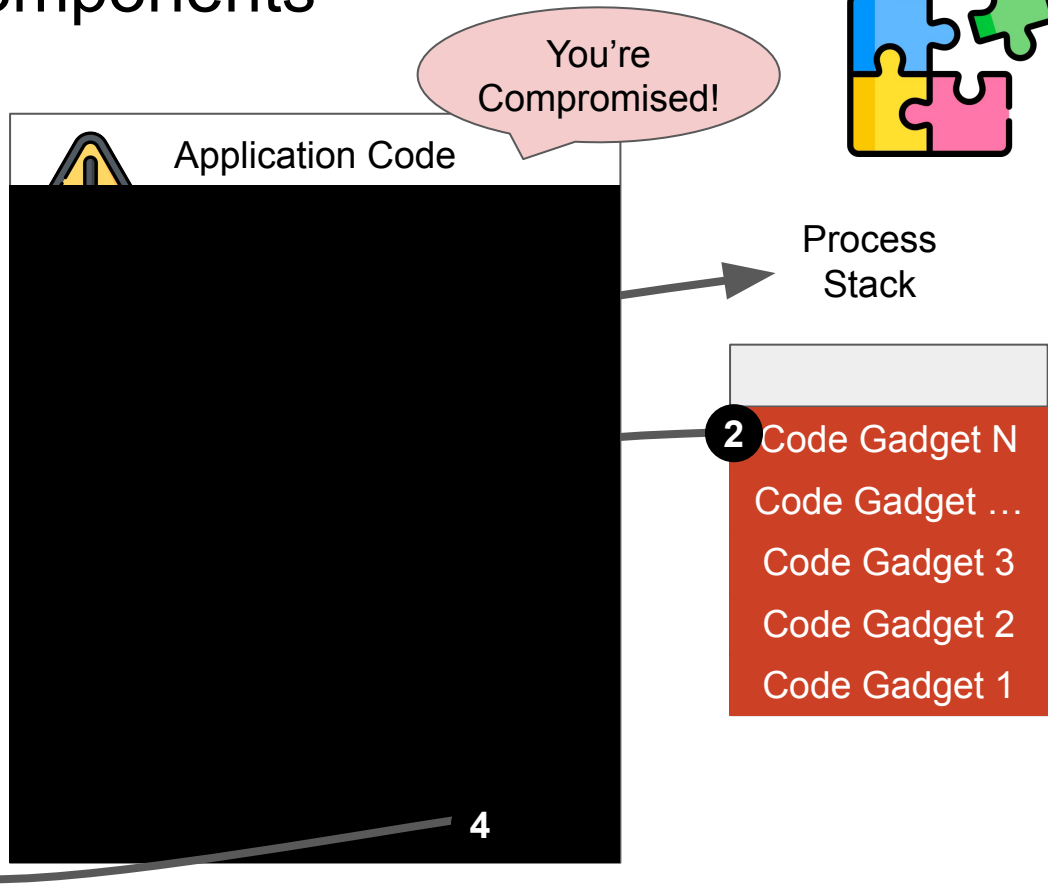


# Outline

- Motivation
- **Background**
  - **Attack-Sensitive Code Components**
  - **How Code Re-Use Attacks Work**
  - **Modern Defense Archetypes**
- Contributions
- Future Work
- Summary

# Attack-Sensitive Code Components

- Code Gadgets
- Pointers (Code & Data)
- System Calls
- Non-Control Data



Process  
Stack

2 Code Gadget N  
Code Gadget ...  
Code Gadget 3  
Code Gadget 2  
Code Gadget 1

4




# How Code Re-Use Attacks Work

## Assumptions / Threat Model

- Memory Vulnerability Exists
- Buffer Overflow
- Dangling Pointer

## Generalized Code Re-Use Attack Procedure

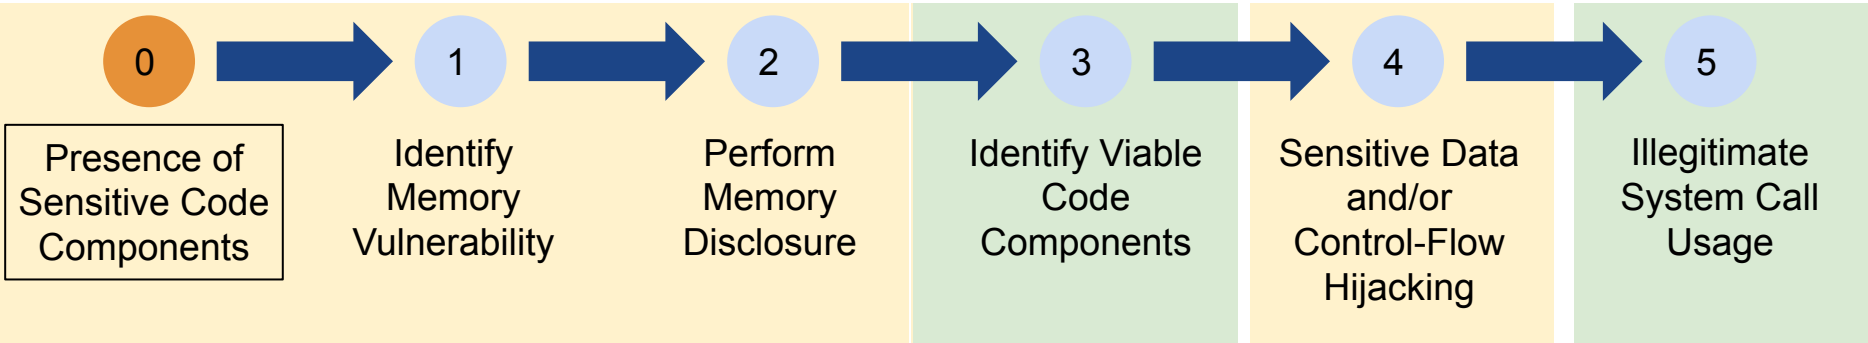


## Code Re-Use Variants

- |   |                                       |
|---|---------------------------------------|
| • Return Oriented Programming (ROP)       | <i>Shacham et al. (CCS 2007)</i>      |
| • Just-In-Time ROP (JIT-ROP)              | <i>Snow et al. (S&amp;P 2013)</i>     |
| • Blind ROP (BROP)                        | <i>Bittau et al. (S&amp;P 2014)</i>   |
| • Control Data & Non-Control Data Attacks | <i>van der Veen et al. (CCS 2017)</i> |

# Modern Defense Archetypes

## Generalized Code Re-Use Attack Procedure




## Generalized Defense Archetype Created

Debloating



Memory Safety



Information Hiding



Re-Randomization



Data Integrity & Control Flow Integrity




System Call Filtering



# Blueprint For Success in Practical Exploit Mitigation Design

## Practical Design Properties


- Low Performance Impact
  - Defenses should achieve **low** performance
- Strong Security Guarantees
  - Defenses should provide **strong** security guarantees
- Scalable Framework
  - Defenses should **minimize** use of additional CPU or memory resources
- Reliable Defense
  - Defenses should **not** break the application runtime



# Blueprint For Success in Practical Exploit Mitigation Design

## Necessary Considerations

- Being Aware of Trade-offs
- Deriving the Essentials
  - What code pieces really needs to be protected?
  - What is the least amount of security coverage needed to still be strong and block attacks?




# Outline

- Motivation
- Background
- **Contributions**
  - MARDU
  - BASTION
- Future Work
- Summary

# Thesis Contributions

This thesis proposal presents two Practical Exploit Mitigation Designs:



- **MARDU**: Efficient and Scalable Code Re-Randomization

- Motivation
- Design
- Implementation
- Evaluation

**SYSTOR '20**  
**DTRAP '22**



- **BASTION**: Context Sensitive System Call Protection

- Motivation
- Preliminary Design

To Be Submitted  
**ASPLOS '23**


# Outline

- Motivation
- Background
- **Contributions**
  - **MARDU**
  - BASTION
- Future Work
- Summary

3


Identify Viable  
Code  
Components

Re-Randomization





# Current randomization techniques are good ...

## Code Randomization



- Address Space Layout Randomization (ASLR)
  - + Light-weight
  - Static code layout
  - One memory disclosure can compromise entire code layout



- Re-Randomization Techniques
  - + Continuous churn makes gadgets hard to find
  - High overhead
  - Rely on predictable thresholds such as
    - Time interval
    - Syscall invocation
    - Call history

Oxymoron  
Coarse-Grained ASLR  
(Memory Page)

*USENIX Security 2014*


Shuffler  
Fine-Grained Runtime  
Re-Randomization  
(Function Granularity)

*USENIX OSDI 2016*



# But They Are Not Practical. Why?

- Users desire **acceptable performance** (<10% avg & worst-case)
- Users desire **strong defenses**
- Users desire **scalability** as more computation is moved to the cloud
  - Have system-wide security coverage including shared libraries
- Users desire **reliable** defense that can be generically deployed
- Achieving all together is **hard**



# Challenges For Practical Re-randomization

- **Performance Challenges**

- Avoiding Useless Overwork: Active randomization wastes CPU cycles in case of “what-if”

- **Security Challenges**

- Code Disclosure: a single leaked pointer allows attacker to obtain code layout of a victim process

- **Scalability Challenges**

- Code Tracking: to support runtime re-randomization tracking and updating of pc-relative code is a necessary and expensive evil
- Stop-the-World: Updating shared code on-the-fly is challenging especially in concurrent access

- **Reliability Challenges**

- Being Generic: Ability for incremental deployment is preferred to completely overhauling a system for isolated protection

# MARDU

## ● MARDU Goal Summary

- |           |             |   |
|-----------|-------------|---|
| ○ Goal 1: | Performance | Low performance impact, but still high entropy  |
| ○ Goal 2: | Security    | Code is secure in runtime                       |
| ○ Goal 3: | Scalability | Code-sharing is possible in MARDU               |
| ○ Goal 4: | Reliability | MARDU supports mixed (un)instrumented libraries |



## MARDU Key Points

- Performance
  - MARDU performs reactive **on-demand** runtime re-randomization
  - No previous randomization scheme is capable of runtime re-randomization **AND** code sharing
- Security / Scalability / Reliability
  - MARDU uses code **trampolines** with PC-relative addressing
  - MARDU can share it's randomized code in a **system-wide** manner



# MARDU

- **MARDU Design**


- MARDU: Design Overview
- Key Concept: How Code Trampolines Work
- Scalability: How MARDU Shares Code
- Performance: Re-randomization Without Stopping the World

- **Implementation**

- **Evaluation**




# MARDU Design Overview

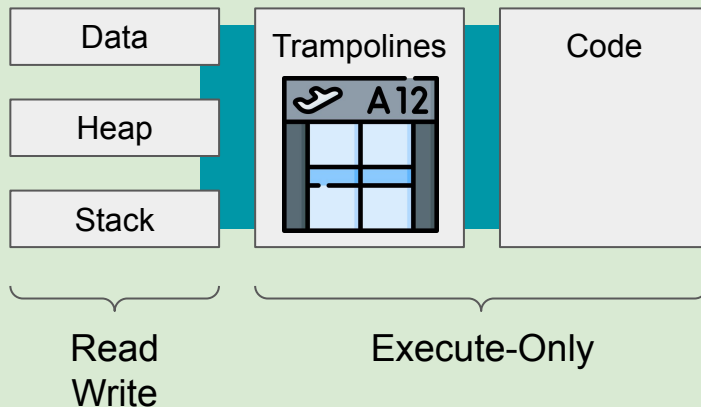


## Compiler Component

- Code Analysis
- Create & Insert MARDU Trampolines
- Generate Fixup Information for Patching



## MARDU Process




**Trampolines** act as:  
**entry gateway** to each function call  
**exit gateway** for each function return

## Kernel Component


- Memory layout & Execute-only Overlay
- On-Demand Re-Randomization & Patching



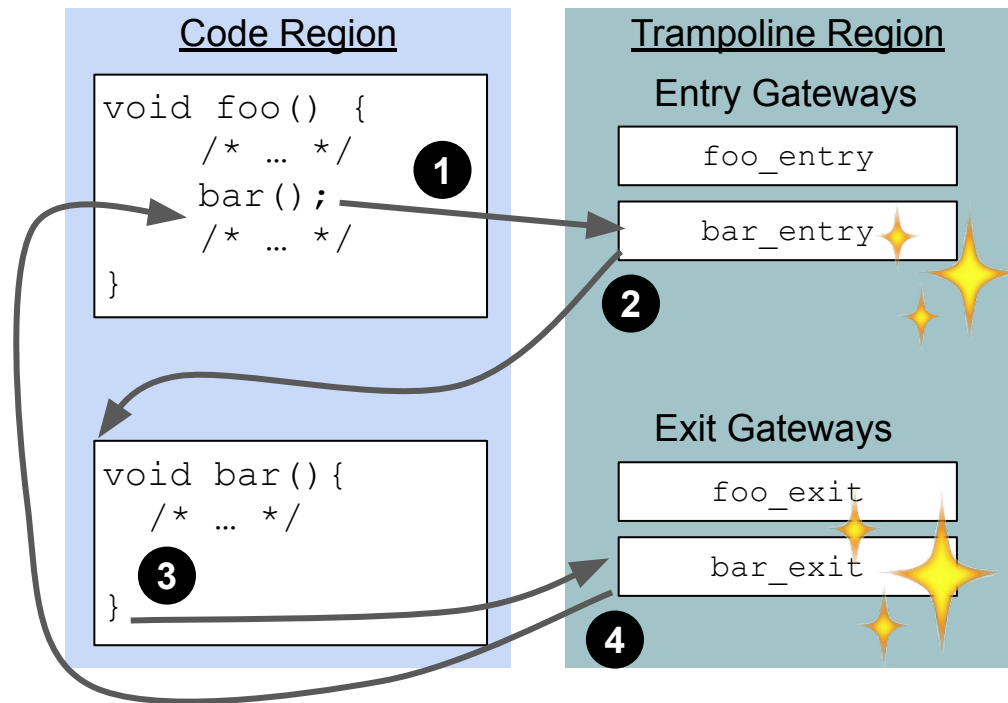
# How Code Trampolines Work



## Traditional Control-Flow






## MARDU Control-Flow

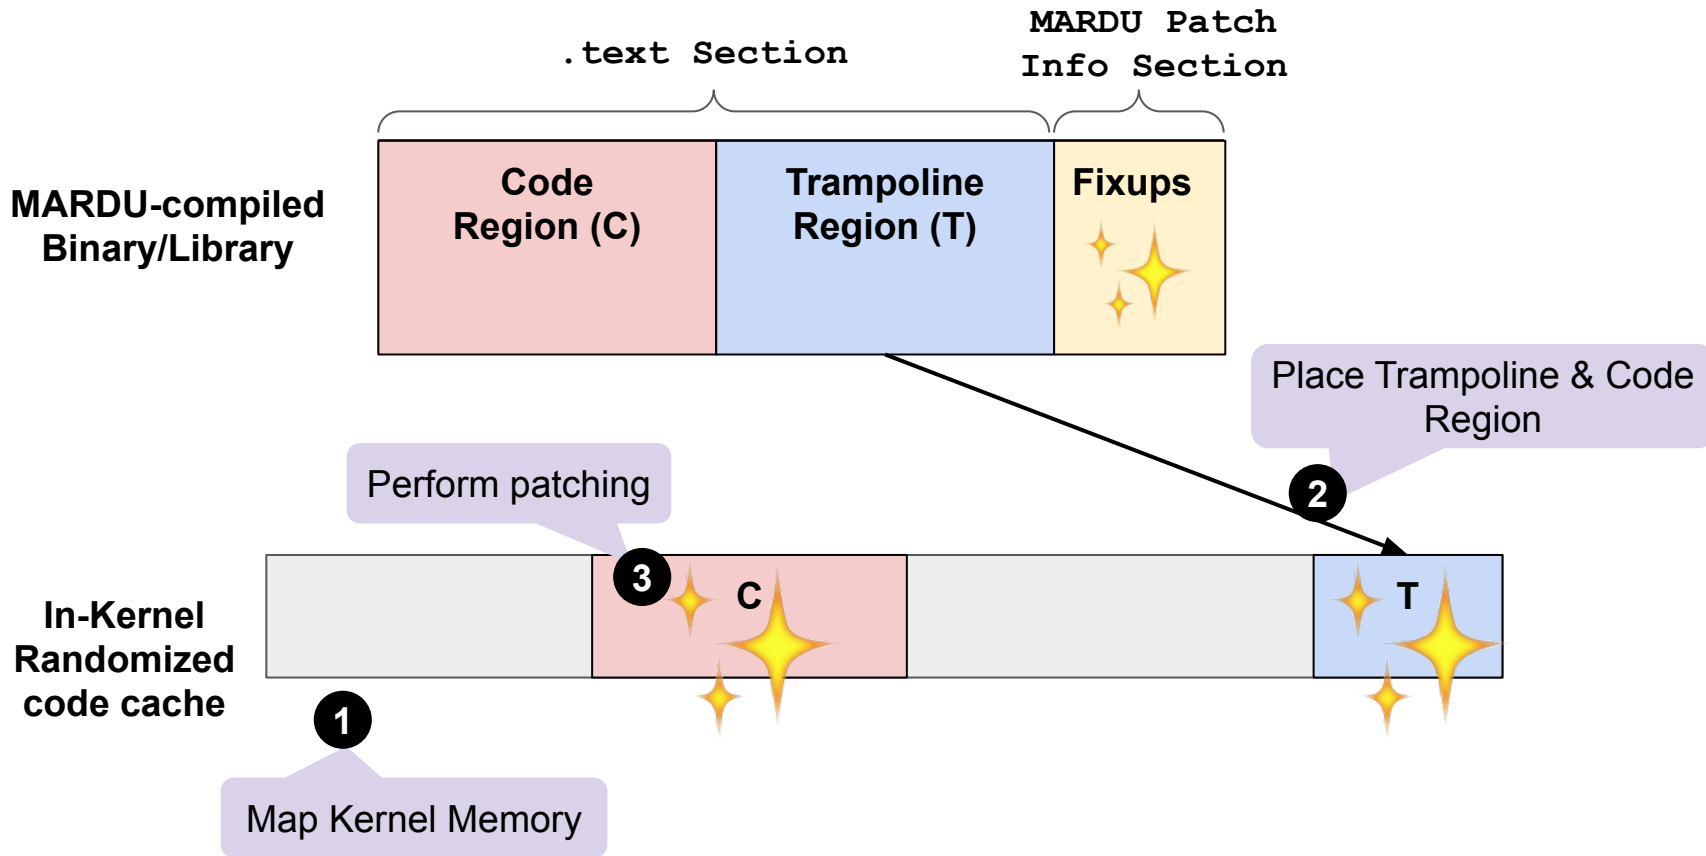



# How Code Trampolines Work

- MARDU runtime always transfers using the trampoline region
- Trampolines leverage immutable code
  - MARDU has a ground truth that will always maintain program semantics
- Simplified Runtime Tracking & Re-Randomization

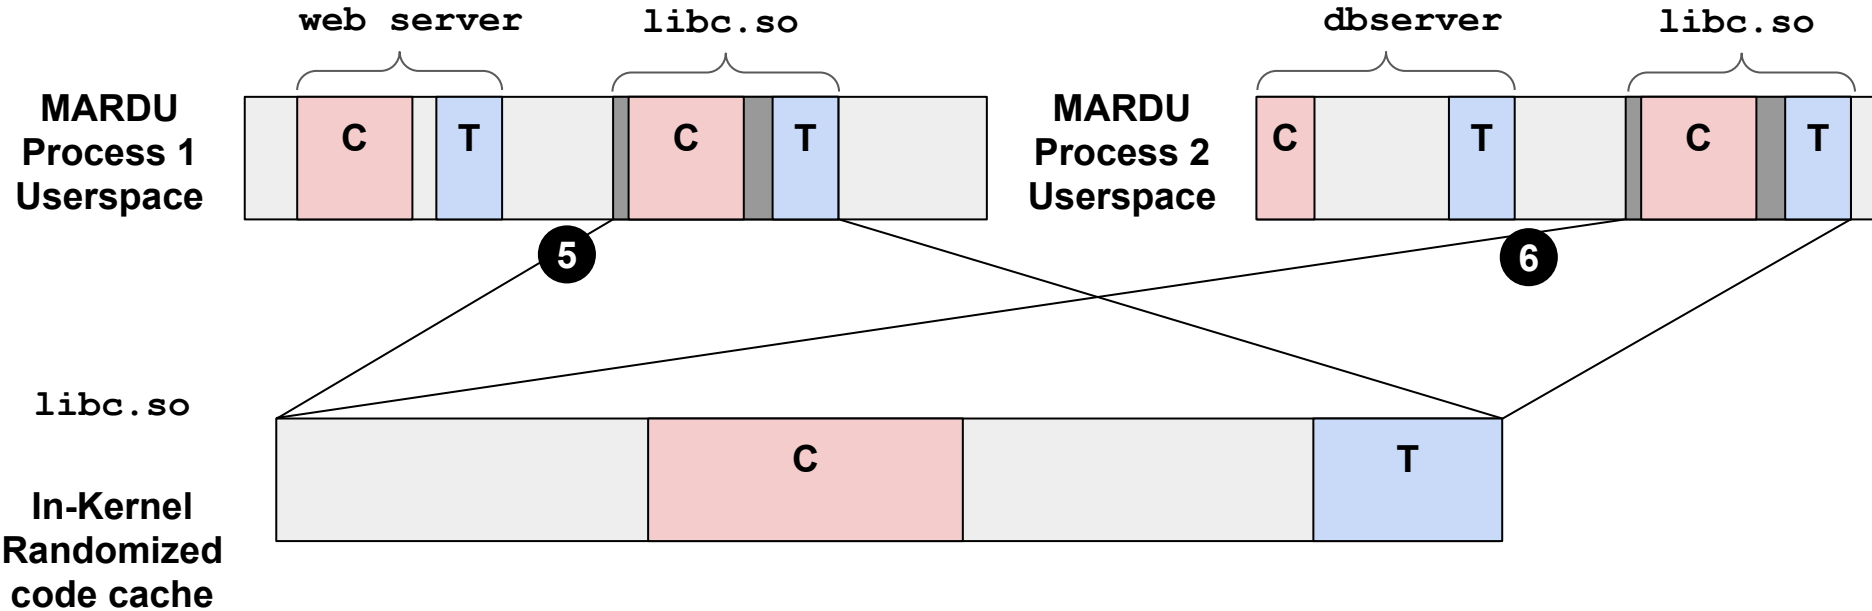


# Example: How MARDU Shares Code

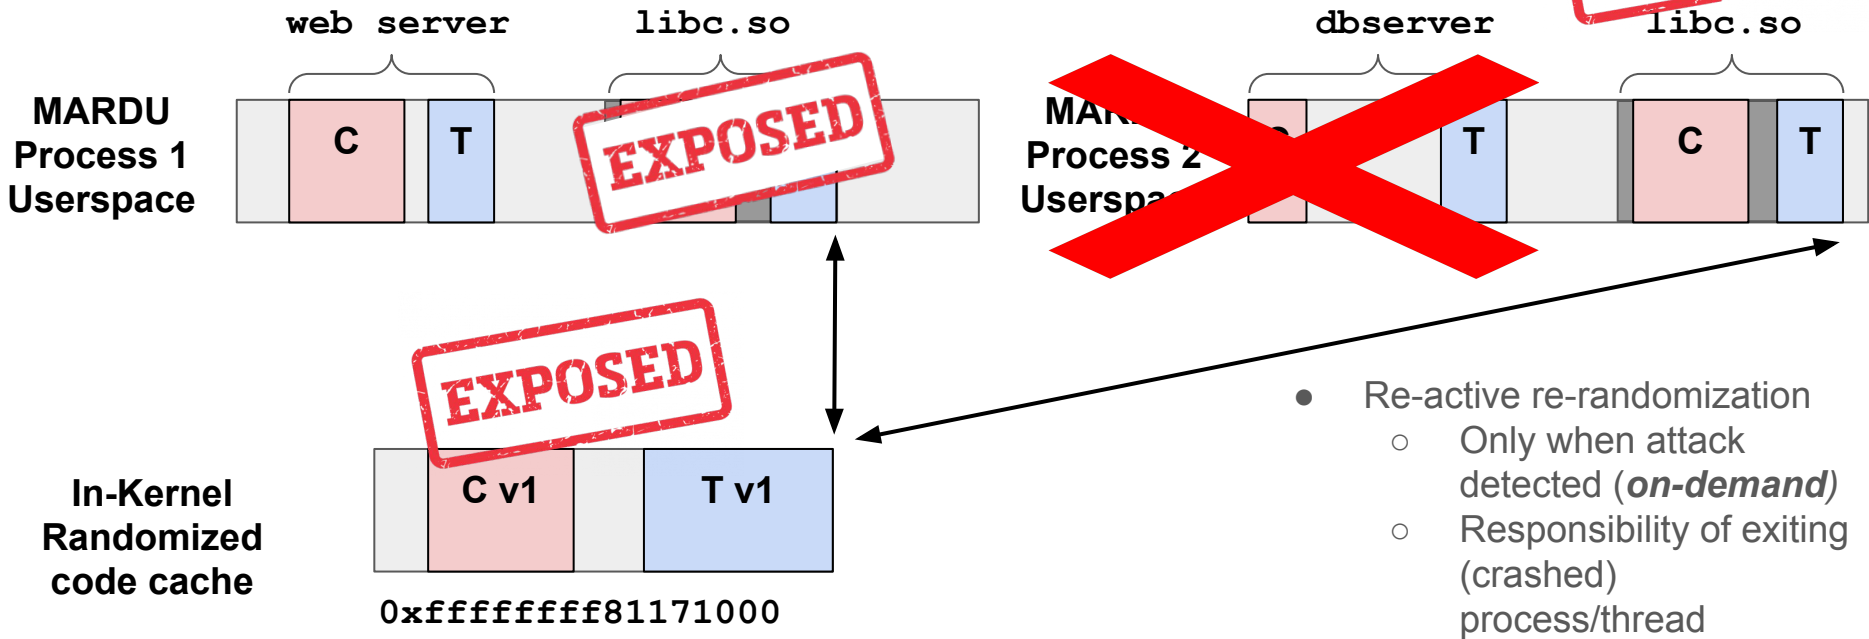




# Example: How MARDU Shares Code




# Re-Randomization Without Stopping the World



# Re-Randomization Without Stopping the World



MARDU  
Process 1  
Userspace

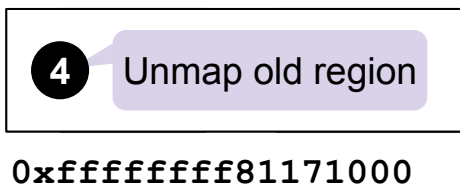


- Gadgets previously deduced are now *stale*
- Randomization is replicated for **ALL ASSOCIATED** shared code of a victim process

2  
Map Code v2 to  
userspace


3  
Map Trampoline v2  
to userspace

In-Kernel  
Randomized  
code cache



# MARDU Implementation

- Working Framework
- Compiler
  - LLVM/Clang 6.0.0
  - 3.5K LOC
- Kernel
  - X86-64 linux 4.17.0
  - 4K LOC
- musl LibC
  - General C library



# MARDU Evaluation

- **Evaluation**

- Security: Popular ROP attacks -> MARDU Wins
- Performance: Compute Bound -> Minimal Runtime Overhead
- Scalability: Concurrent WebServer-> Negligible Runtime Overhead and Scalability

- 1) How much performance overhead does MARDU impose?
- 2) How scalable is MARDU in terms of memory savings?

# Experimental Setup and Applications

- Experimental Setup

- All programs compiled with MARDU LLVM compiler and `-O2 -fPIC` optimization flags
- Platform:
  - 24-core (48-Hardware thread) machine with two Intel Xeon Silver 4116 CPUs (2.10 GHz)
  - 128 GB DRAM


- Applications

- SPEC CPU 2006 (All C applications)
- NGINX web server


# How MARDU Performs



## CPU Intensive Benchmark SPEC CPU 2006



## Concurrent Web Server NGINX




NGINX Average Degradation: 4.4%

# How MARDU Scales



- One-time upfront cost to instrument
- **Native** shared library total size: **787 MB**
- **Traditional non-sharing** would incur **8.8 GB** overhead
  - Library replicated & allocated new memory per process
- MARDU incurs **1.3 GB** memory usage (over **7.5 GB** memory savings)
- Biggest memory savings come from
  - **libc.so** saves **0.78 GB** memory
  - **libm.so**, saves **0.26 GB** memory





# MARDU Summary & Limitations

## Summary

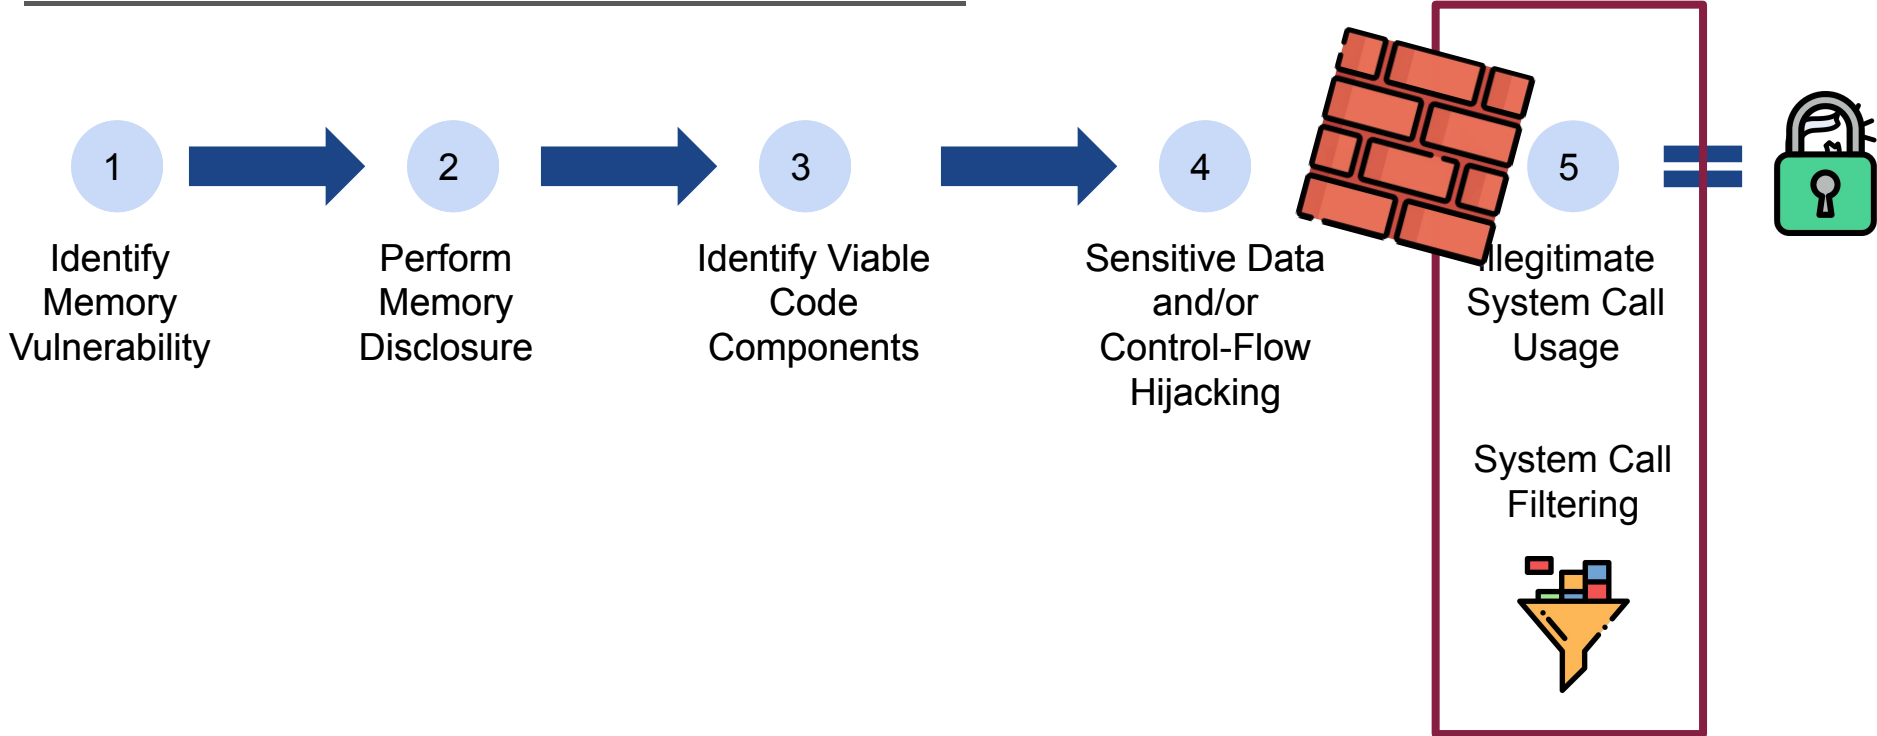
- MARDU is Practical
- 5.5% average performance overhead
- Security blocks all ROP attack variants
- First randomization scheme capable of runtime re-randomization **with** code sharing
- Does not Stop-the-World for runtime re-randomization

## Limitations

- Re-Randomization is a **probabilistic** defense
- Performance could be improved, worst case performance is 18.3%
- **Non-Control Data Attacks** are not covered by MARDU


# Modern Security

## Generalized Code Re-Use Attack Procedure



# Outline

- Motivation
- Background
- **Contributions**
  - MARDU
  - **BASTION**
- Future Work
- Summary



5

Illegitimate  
System Call  
Usage

System Call  
Filtering




# Outline

- Motivation
- Background
- **Contributions**
  - MARDU
  - **BASTION**
    - Insights & Motivation
    - Design
    - Preliminary Evaluation
- Future Work
- Summary


5

Illegitimate  
System Call  
Usage

System Call  
Filtering




# Insights of System Call Usage in Code Re-Use



## For **Applications** (Legitimate Usage)


- Provide an API interface between applications and host OS kernel
- Provide **numerous services**
- **359** system calls currently implemented (Linux v5.17.1)
- Many system calls are **non-sensitive**
- System Calls are **scarcely** used in practice



## For **Attacks** (Malicious Usage)

- Provide an API interface to further gain a foothold on victim host
- Some system calls are **security-critical** especially:
  - `execve`
  - `mmap`
  - `mprotect`

# Example 1: Duality of System Call Usage



## NGINX Web Server

### Legitimate Uses

- **execve()** used to update server in place during runtime

### Attacker Uses

- **execve()** can launch an attacker binary by reaching system call and corrupting `ctx->path`, `ctx->argv`


```
// nginx/src/os/unix/nginx_process.c
static void ngx_execute_proc(ngx_cycle_t *cycle, void *data) {

    ngx_exec_ctx_t *ctx = data;

    if( execve( ctx->path, ctx->argv, ctx->envp ) == -1 ) {

        ngx_log_error(NGX_LOG_ALERT, cycle->log,
                      ngx_errno, "execve() failed");
    }
    exit(1);
}
```

# Summary of System Call Usage



## Code Re-Use Attacks

- Attacks greatly vary in approach, complexity, and end goals
- Attacks desire to **reach** and **use system calls** in needed way

## Consequences:

- System calls could be considered a **critical lynchpin** in attack completion
- Protecting system calls can block many system-call based attacks
- **Strong contexts** are needed to adequately protect system calls

## Bottom Line:

- System calls deserve more attention!

# Current Practices in System Call Filtering

## What is Current System Call Filtering Doing?

- Analysis to derive Allow/Block list of system calls
- Static Argument Constraints

### **seccomp: System Call Whitelisting<sup>1</sup>**

- Defacto **manual** system call policy mechanism implemented in Linux
- Needs **new separate policy** for every application, & every configuration

### **sysfilter: Automated Policy Generation<sup>2</sup>**


- Introduction of automation from analysis framework to generate appropriate `seccomp` filter for a target application
- reduces burden on administrator, not stronger security

<sup>1</sup> The kernel development community. Seccomp BPF (SECure COMPUting with filters), 2015

<sup>2</sup> Nicholas DeMarinis et al. sysfilter: Automated system call filtering for commodity software. (*RAID 2020*)



# BASTION: Our Solution



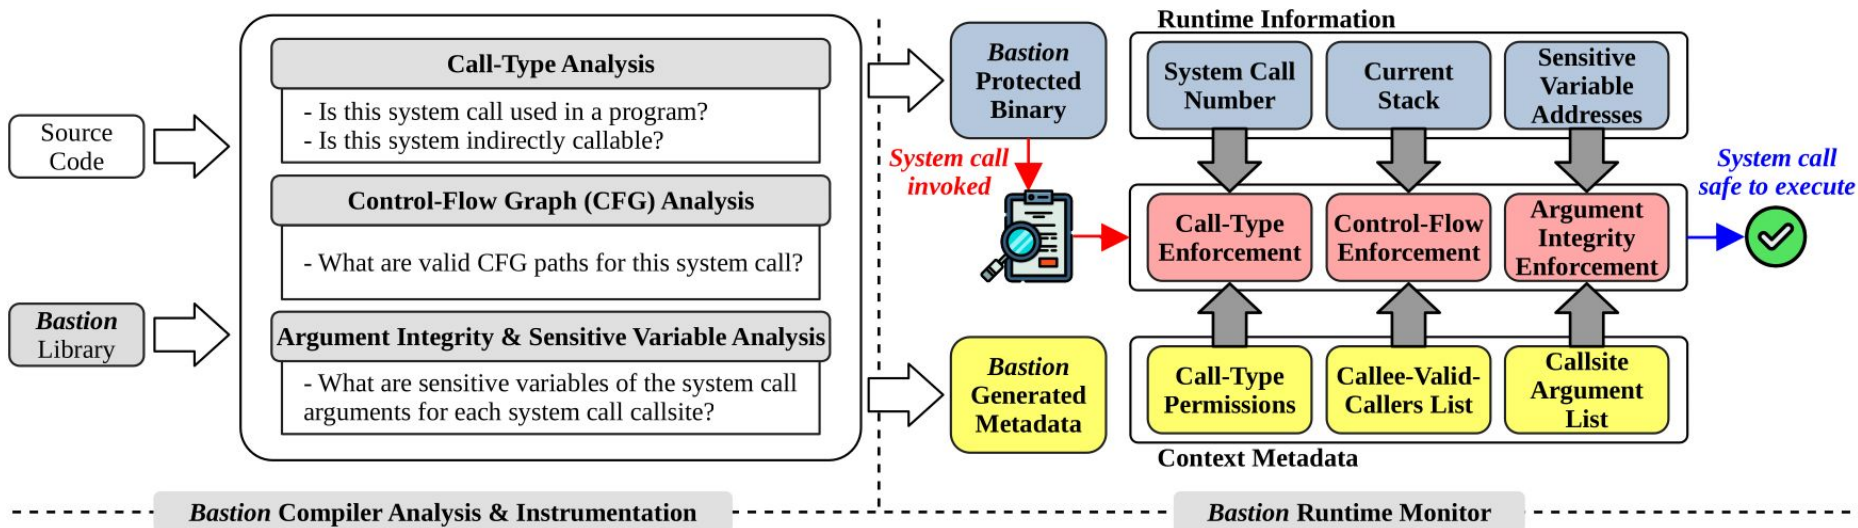

- *Surround* system calls with **three** tight, **specialized** contexts

## BASTION System Call Contexts


- **Call-Type** Context
  - Is this system call used in the program?
  - How is this system call invoked by the program? (directly, indirectly, or both)
- **Control-Flow** Context
  - What are the valid Control Flow Graph (CFG) paths that reach this system call callsite?
- **Argument Integrity** Context
  - What are the sensitive variables used as system call arguments and dependent variables for each system call callsite?

# BASTION: Our Solution

## BASTION Framework Overview



# BASTION: Preliminary Evaluation




## Evaluation

**Application:** NGINX Web Server

## Results


- Negligible Overhead
- BASTION Argument Integrity is *significantly cheaper* than DFI by being *specialized*



# Outline

- Motivation
- Background
- Contributions
- **Future Work**
- Summary

# BASTION Future Work



## Current Status:

- Bare-Bones Prototype of BASTION in place


## Ongoing Research

1. Complete BASTION System Call Defense Framework
2. Comprehensive **Security** Study of BASTION
3. Comprehensive **Performance** Study of BASTION

## Questions To Answer:

- Can BASTION adequately protect against prominent Code Re-Use Attack Vectors?
- Is BASTION just as performant for a diverse set of real-world applications as shown in our preliminary evaluation?

# BASTION Future Work



## Proposed Security Study:

- Real-world CVEs
- Address Oblivious Code Re-Use (AOCR)
- NEWTON: Dynamic Gadget Discovery Framework

*NDSS '17*

*CCS '17*

## Proposed Performance Study:

- NGINX Web Server
- SQLite SQL Database Engine

## Other real-world application candidates


- vsftpd FTP server

# Outline

- Motivation
- Background
- Contributions
- Future Work
- **Summary**

# Summary of Research

- Defense techniques to fight *code re-use*
- Exploit Mitigation Mechanisms that are *practical*



## MARDU Key Contributions

- Enabled capability of **code sharing** for randomized code
- Efficient, **system-wide** (re-)randomization that is **on-demand**



## BASTION Key Contributions

- Stop code re-use attack chain at system call usage step
- **Three new contexts** to protect system calls
- **Narrow** design only enforcing integrity on system call relevant components



# Summary of Research

Publications:

## BASTION

**BASTION: Protect the System Call, Protect (most of) the World**

- To be completed and submitted *ASPLOS '23*

## MARDU

**MARDU: Efficient and Scalable Code Re-randomization**

**Christopher Jelesnianski, Jinwoo Yom, Changwoo Min, and Yeongjin Jang (*SYSTOR'20*)**

**Securely Sharing Randomized Code that Flies**

**Christopher Jelesnianski, Jinwoo Yom, Changwoo Min, and Yeongjin Jang (*DTRAP'22*)**

# Summary of Research

## Other Publications:

### Security

#### **Tightly Seal Your Sensitive Pointers with PACTight**

Mohannad Ismail, Andrew Quach, **Christopher Jelesnianski**, Yeongjin Jang and Changwoo Min (*To appear USENIX Security'22*)

- Data protection utilizing ARM Pointer Authentication (PA) security primitive

#### **VIP: Safeguard Value Invariant Property for Thwarting Critical Memory Corruption Attacks**

Mohannad Ismail, Jinwoo Yom, **Christopher Jelesnianski**, Yeongjin Jang and Changwoo Min (*CCS'21*)

- Value Integrity for security-relevant data types

### Compilers & System Software

#### **Breaking the Boundaries in Heterogeneous-ISA Datacenters**

Antonio Barbalace, Robert Lyerly, **Christopher Jelesnianski**, Anthony Carno, Ho-Ren Chuang, Vincent Legout, and Binoy Ravindran (*ASPLOS'17*)

#### **Operating system process and thread migration in heterogeneous platforms**

Robert Lyerly, Antonio Barbalace, **Christopher Jelesnianski**, Vincent Legout, Anthony Carno, and Binoy Ravindran (*MaRS'16*)

#### **Popcorn: Bridging the programmability gap in heterogeneous-isa platforms**

Antonio Barbalace, Marina Sadini, Saif Ansary, **Christopher Jelesnianski**, Akshay Ravichandran, Cagil Kendir, Alastair Murray, and Binoy Ravindran (*Eurosys'15*)



# **Practical Exploit Mitigation Against Code Re-Use Attacks and System Call Abuse**

Christopher Jelesnianski

Thank you!

Questions?

