


PRISM: Optimizing Key-Value Store for Modern Heterogeneous Storage Devices

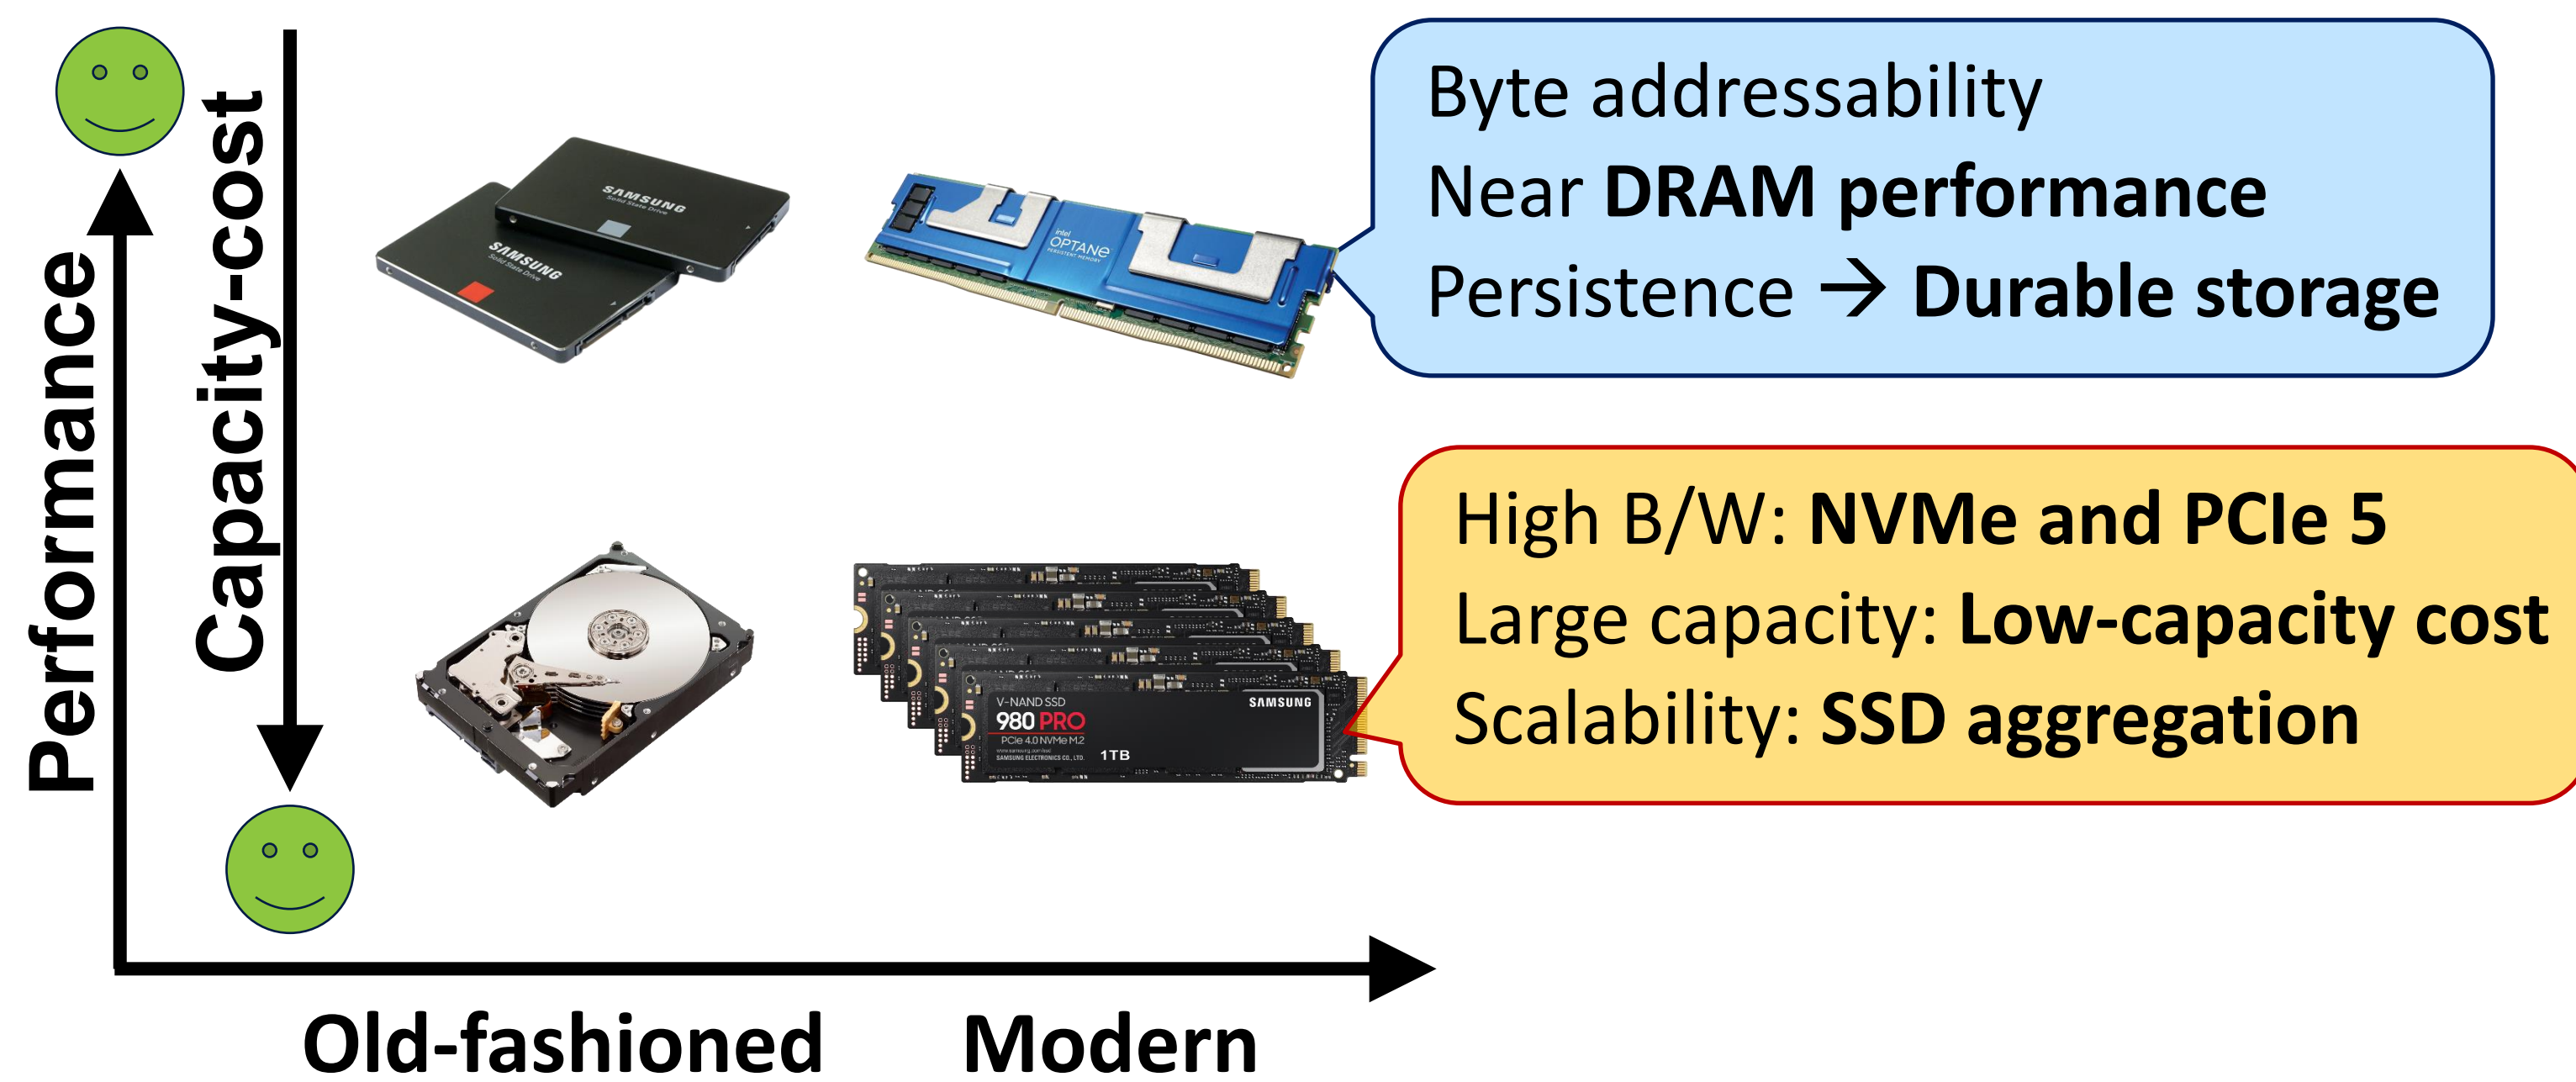
Yongju Song, Wook-Hee Kim¹, Sumit Kumar Monga², Changwoo Min², and Young Ik Eom
Sungkyunkwan University, ¹Konkuk University, ²Virginia Tech



Scan for Full Paper

Q. How should we design a Heterogeneous Storage System in a Modern Storage Landscape?

1. Modern Heterogeneous Storage Devices




2. Evolution of Storage Heterogeneity

- No clear separation between performance/capacity layers.
 - ✓ “The Storage Hierarchy is *Becoming a Jungle*.” [CIDR’21, Dong Xie]
 - ✓ “The Storage Hierarchy is *Not a Hierarchy*.” [FAST’21, Remzi H. Arpaci-Dusseau]

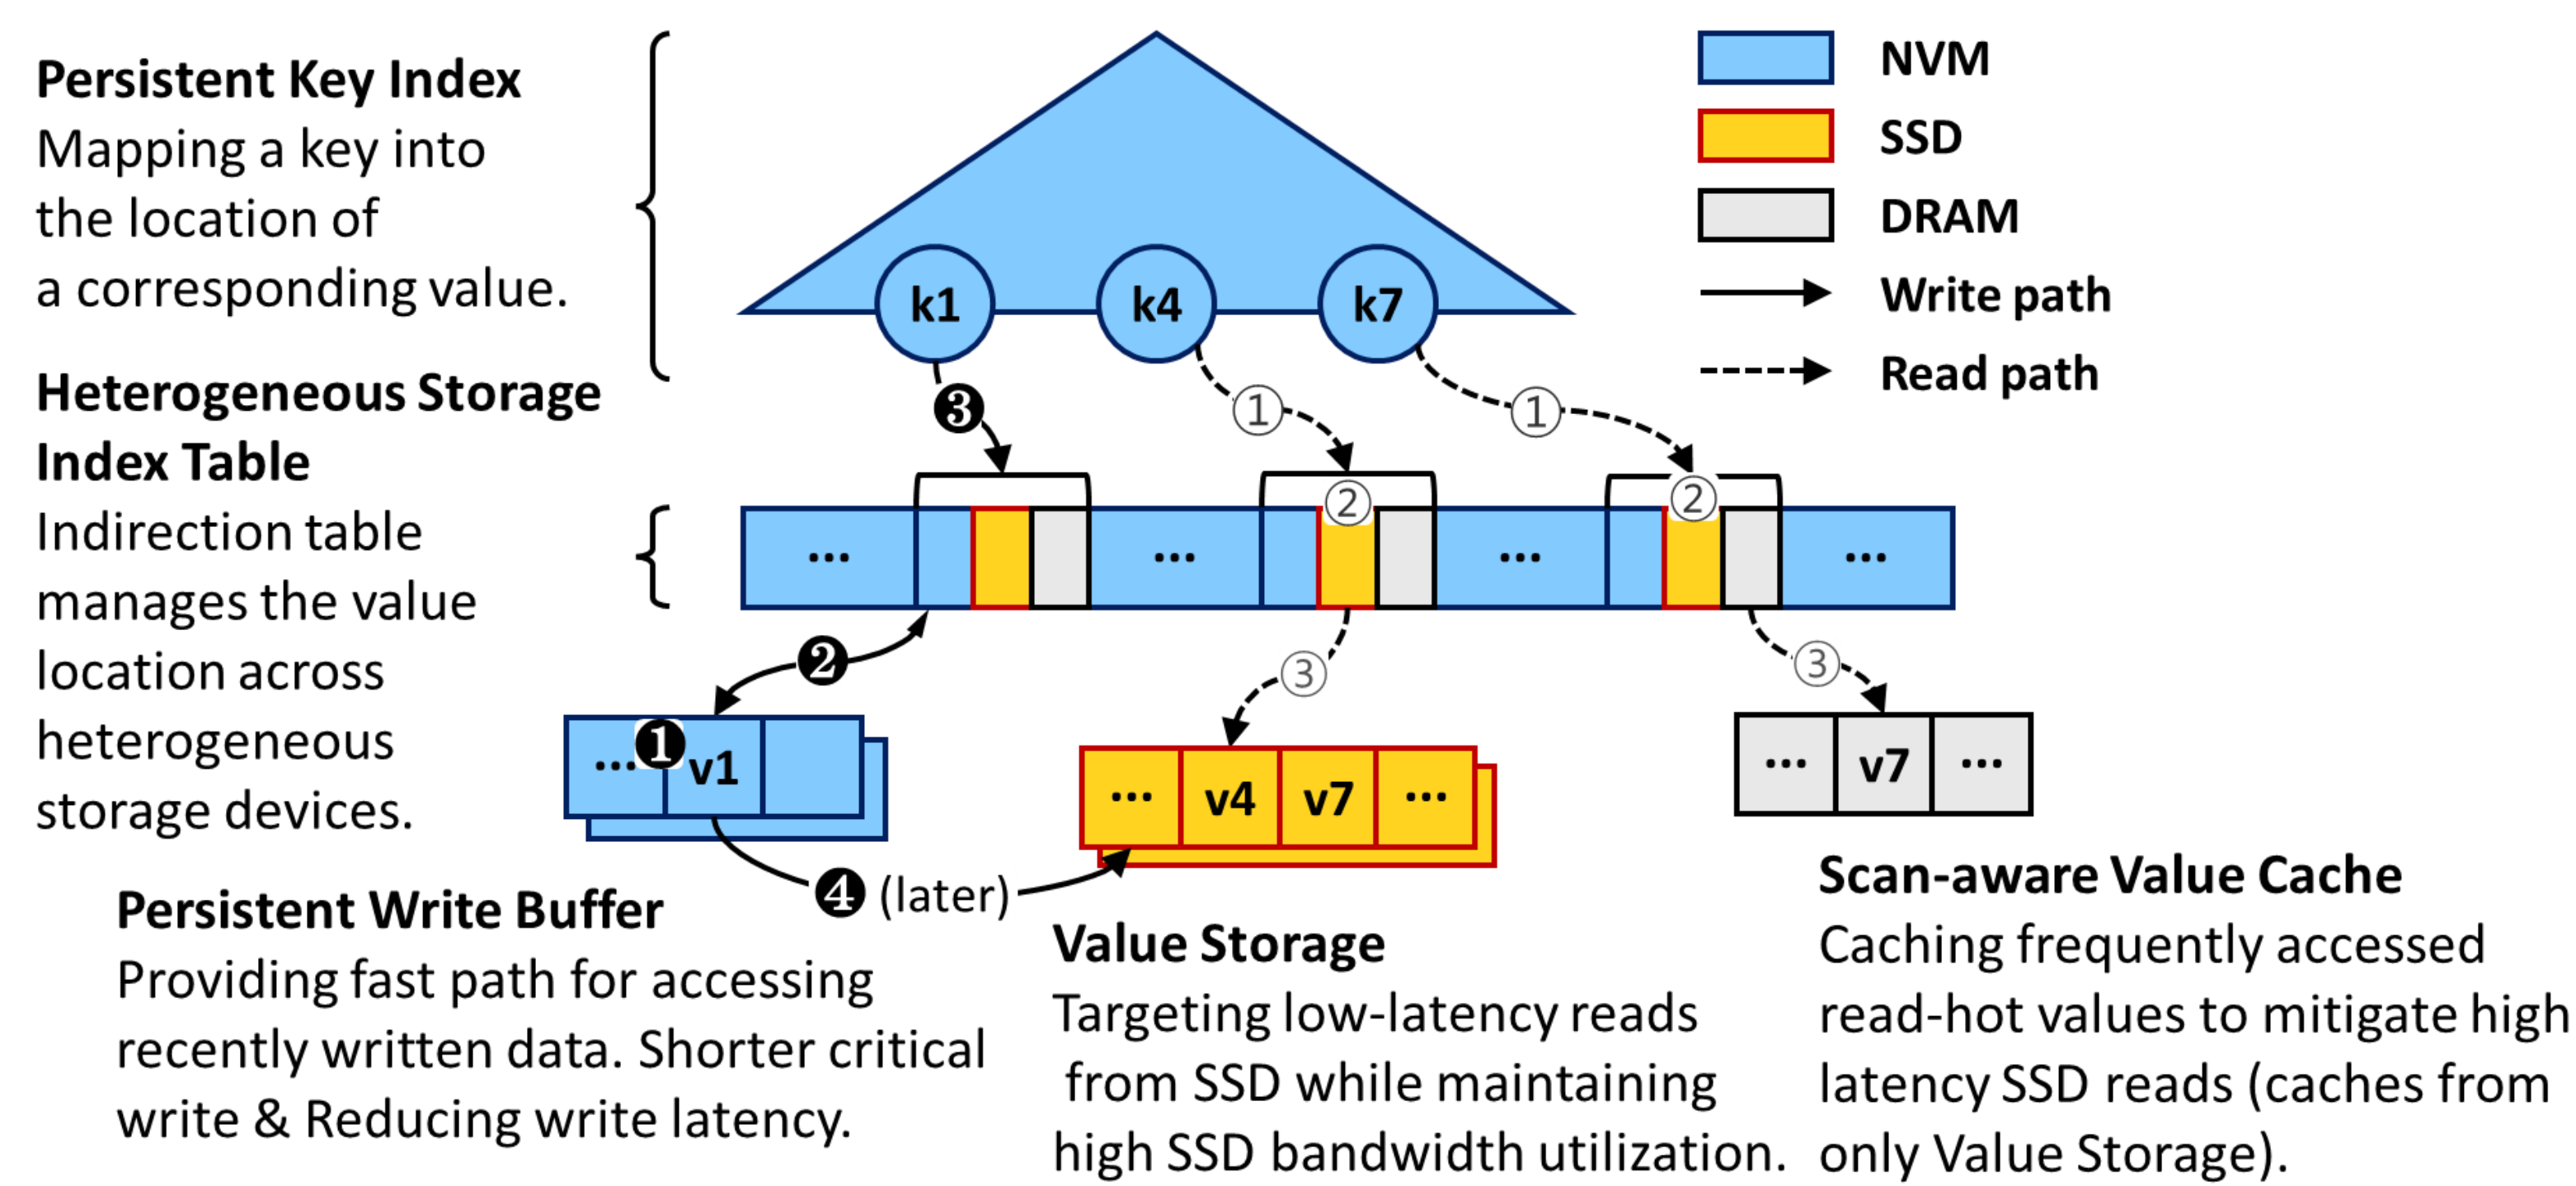
Model	GB	\$/TB	Performance (μs, GB/s)				Endurance Warranty (PBW)
			Read Latency	Write Latency	Read BW	Write BW	
SK Hynix DRAM	16	5,427	0.08	0.08	15	15	∞
Intel Optane DCPMM	128	4,096	0.30	0.09	6.8	1.9	292
Intel Optane 905P PCIe 3	960	400	10	10	2.6	2.2	17.52
Samsung 980 Pro PCIe 4	1,024	100	50	20	7	5	0.6
Samsung 980 PCIe 3	1,024	70	60	20	3.5	3	0.6

3. Managing Storage Heterogeneity Today

- Placing hot data on NVM
 - ✓ System can leverage the low latency of NVM but *suffer from NVM’s limited bandwidth*.
- Traversing data layer by layer for handling read requests
 - ✓ Inefficient traversal leads to *wasting CPU cycles*.
 - ✓ Overall performance may be *bounded to the device with the lowest performance*.

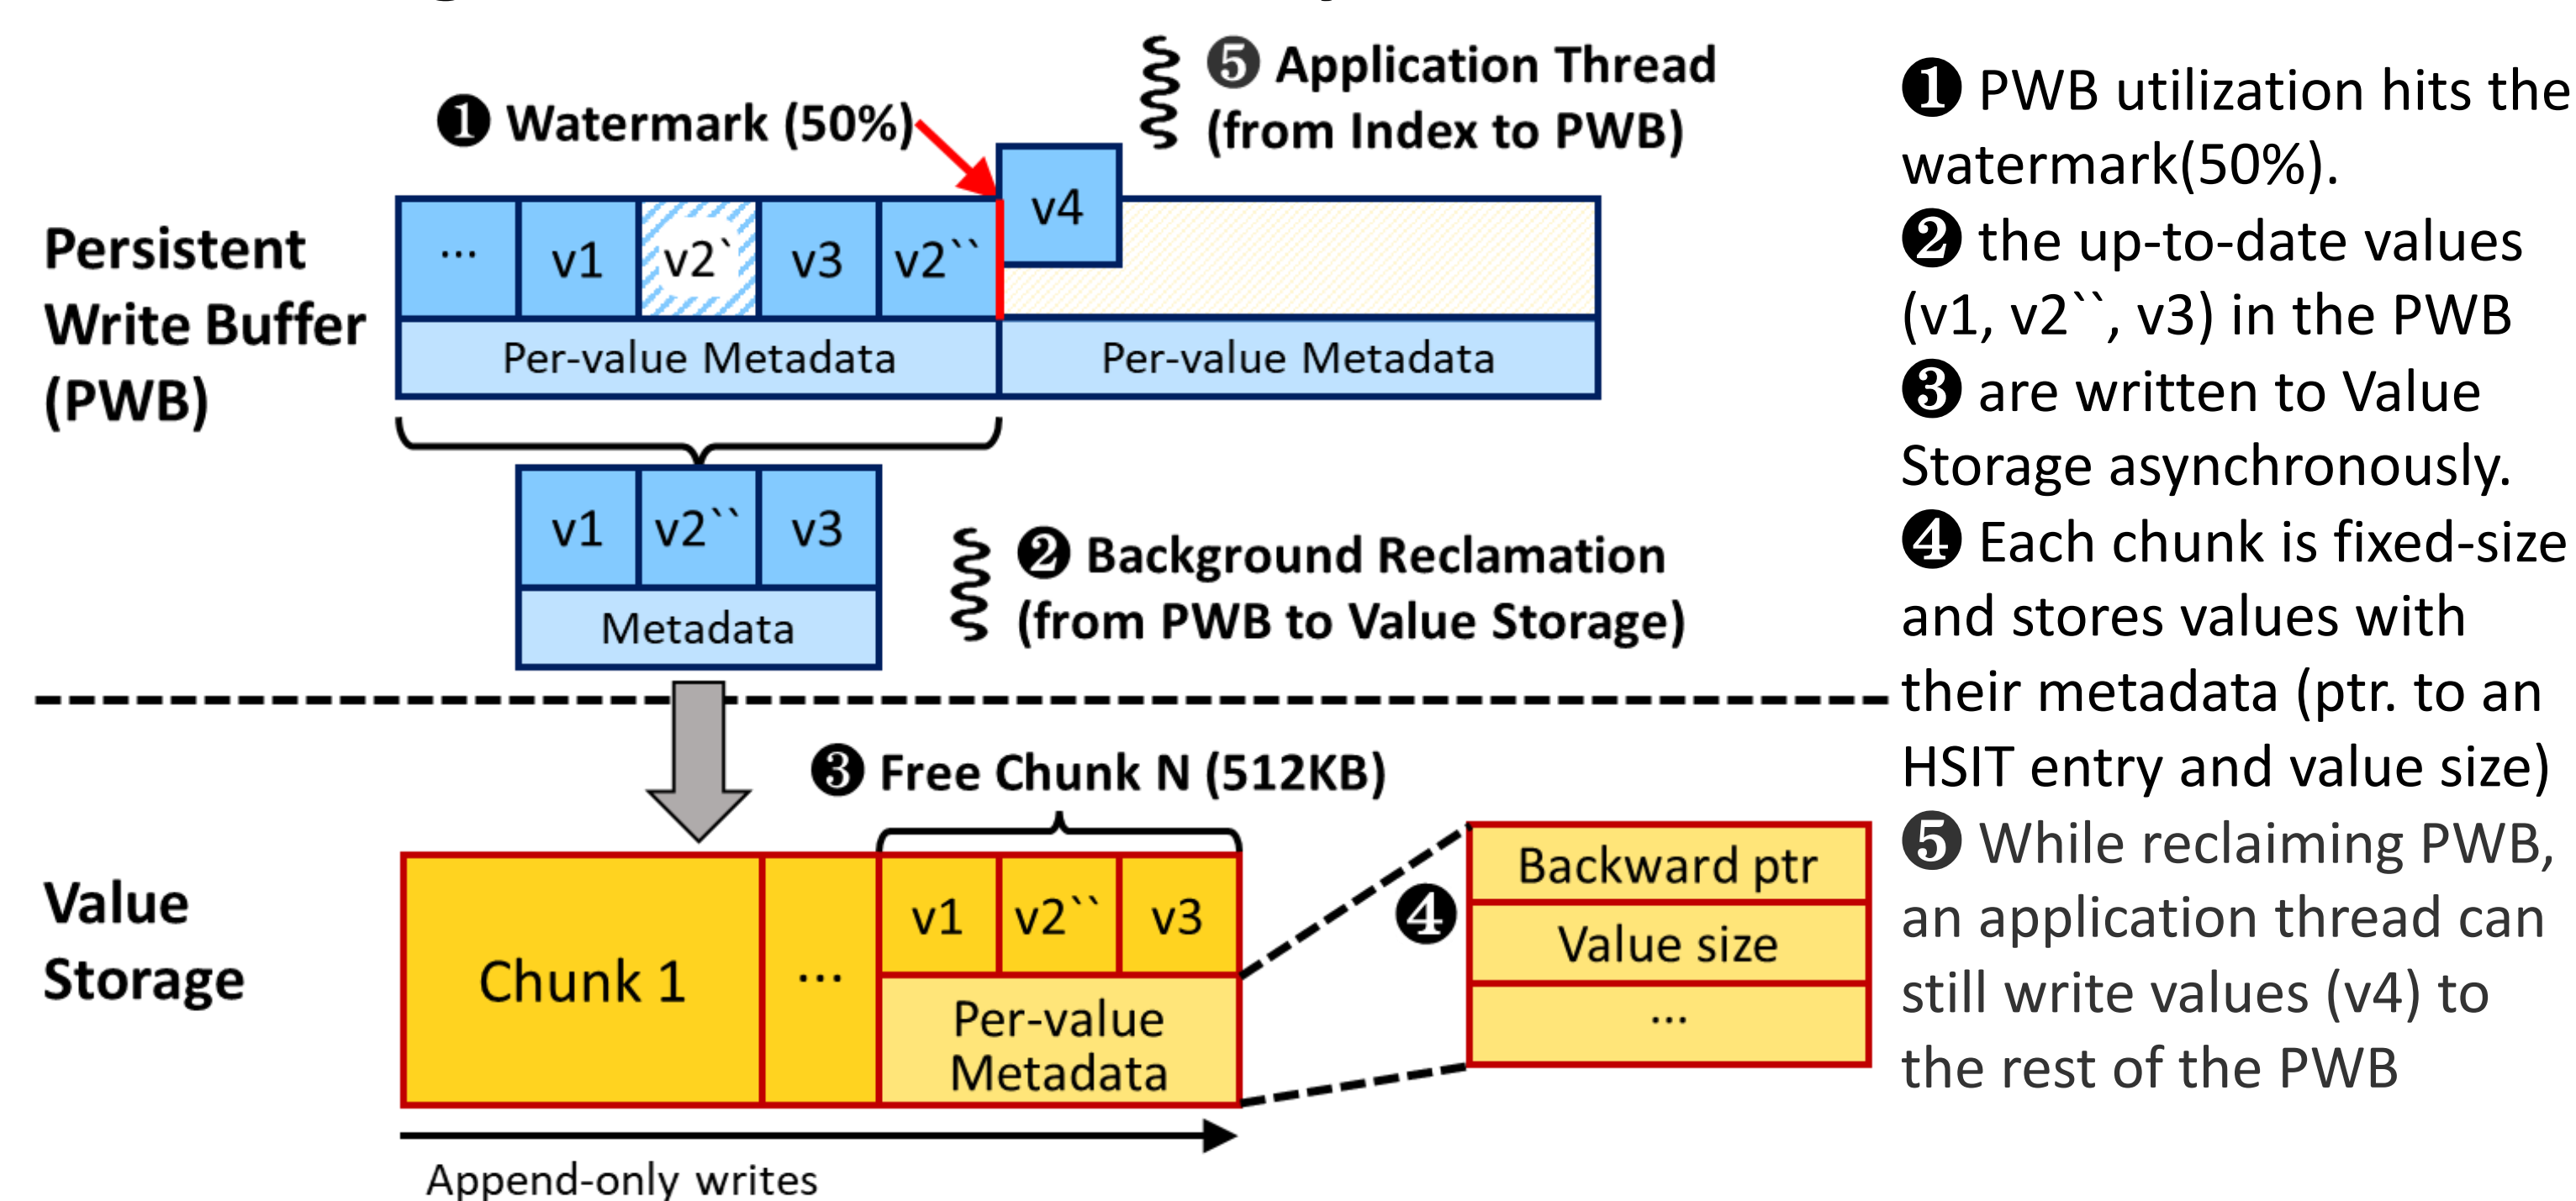


4. Design Overview of PRISM



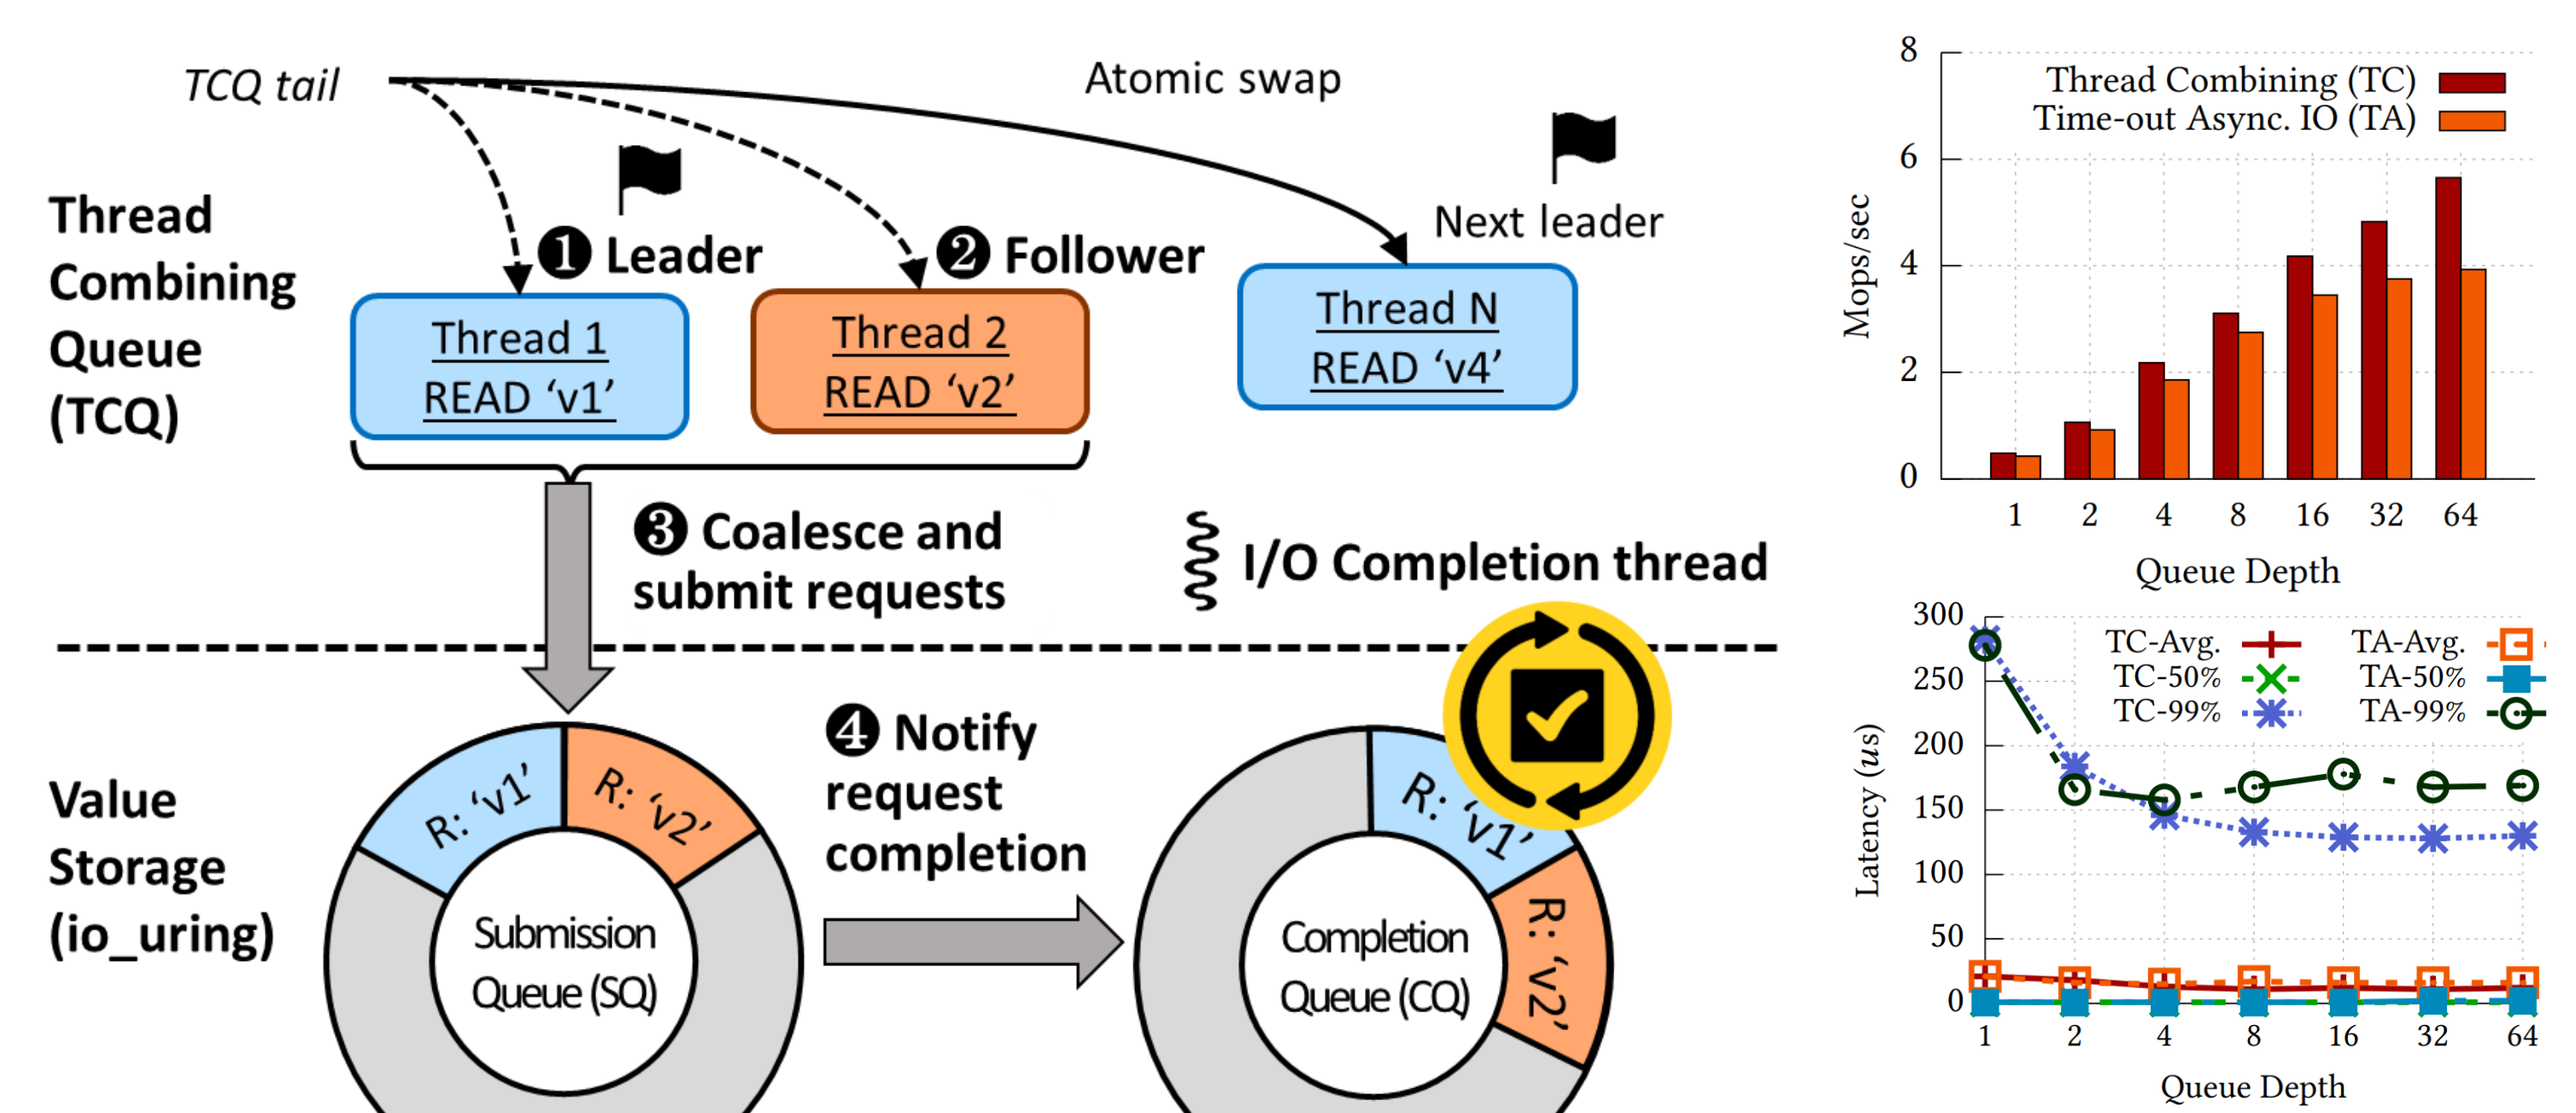
5. Asynchronous Bandwidth-Optimized Write

- Background reclamation : Preventing App. Thread from blocking
- Asynchronous IO batching: Achieving high bandwidth of SSDs
- Allocating a free chunk is the only CS: Concurrent Writes




6. Opportunistic Thread Combining Read

- Dynamically determining the right IO batch size (High B/W & Low Lat.)
 - ✓ Leader dynamically coalesces read req. of followers and submit
 - ✓ Conditions for submitting requests: (prerequisite: Value Storage is idle).
 1. No more followers or 2. When leader thread reaches limit QD.




7. Cross-media Crash Consistency

- Crash consistent update of values on PWB & Value Storage w/ HSIT.
 - ✓ First writes a value with a backward pointer (③, ③), and then it updates its forward pointer (④, ④), invalidating the old forward pointer (②, ②).
- Efficiently guarantees cross-media crash consistency with our pointer update protocol and append-only PWB write policy.



8. Evaluation



- Other interesting evaluations and in-depth analysis
 - ✓ Multicore Scalability, Data Skewness, and Write Amplification
 - ✓ Performance Impact from various system configurations
 - ✓ More details for understanding PRISM performance
- Available in GitHub at <https://github.com/cosmos-jigu/prism>